

Measuring Heterogeneity in Consumption Behavior: Micro Drivers and Macro Implications

July 18, 2018

Edmund Crawley¹

JHU and Danmarks Nationalbank

Andreas Kuchler²

Danmarks Nationalbank

Abstract

This paper uses a novel econometric method to estimate households' consumption responses to permanent and transitory shocks to income. Our method builds upon Blundell, Pistaferri, and Preston (2008), using techniques from Crawley (2018) to correct for the time aggregation problem. Applying our method to a confidential panel dataset covering the entire Danish population, we find heterogeneity in consumption behavior across dimensions of particular importance for monetary policy. Households who stand to lose from an interest rate hike have significantly larger consumption responses than those who stand to gain. Overall we estimate this interest rate exposure channel to be of an order of magnitude larger than the intertemporal substitution channel that dominates in representative agent New Keynesian models.

Keywords Uncertainty, Consumption Dynamics, MPC

JEL codes D12, D31, D91, E21

¹Crawley: Department of Economics, Johns Hopkins University, ecrawle2@jhu.edu ²Kuchler: Danmarks Nationalbank, aku@nationalbanken.dk

Thanks to Danmarks Nationalbank for providing financial support as well as invaluable access to their data.

1 Introduction

This paper makes two contributions to the literature on household consumption behavior and its implications for the macroeconomy. First, we develop a new methodology to estimate the consumption responses to income shocks using covariance restrictions on household level panel data. Second, using administrative data covering the entire Danish population, we use the method to quantify the importance of heterogeneity in monetary policy transmission. For example, we estimate the interest rate exposure channel, which is mute in representative agent models, to be of an order of magnitude larger than the intertemporal substitution channel that typically dominates in New Keynesian models.

Our method is closely related to Blundell, Pistaferri, and Preston (2008) (henceforth BPP) who estimate ‘insurance’ parameters to transitory and permanent shocks to income using covariance restrictions on a panel of household income and consumption. The BPP methodology (or those closely related to it) has become standard in the literature¹ despite the fact that the estimate of the consumption response to transitory income shocks using this method is often an order of magnitude lower than that implied by the natural experiment literature. Crawley (2018) shows one way in which this methodology is seriously flawed: a failure to account for the time aggregation problem of Working (1960). He shows that in a simulated dataset where households follow the permanent income hypothesis, that is they respond one-for-one to shocks to permanent income but not at all to transitory income shocks, the estimate for the consumption response to transitory income shocks using the BPP methodology is -0.6 (a 10% transitory increase to income is estimated to result in a 6% decrease in consumption, where the true change is 0%). Our methodology not only correctly accounts for the time aggregation problem but is also robust to generic short term dynamics in both income and consumption. Our empirical results are also broadly in line with the long literature that uses natural experiments to measure marginal propensities to consume out of transitory income. The popularity of the BPP methodology, despite its flaws, shows the high demand for a method of extracting consumption behavior from panel data on income and consumption. We believe our method, or slight variants of it, provides a robust alternative to BPP and has a wide array of potential applications, particularly when applied to the growing number of large panels such as the Danish administrative data used in this paper, or those available from financial aggregation platforms.²

The need for better methods and data to measure consumption behavior at the household level has grown with the increasing recognition that household heterogeneity may play a key role in macroeconomic dynamics. Kaplan and Violante (2018) provide a nice overview of the theoretical literature incorporating household heterogeneity into models of economic fluctuations. Computational and methodological limitations, along with early work by Krusell and Smith (1998) showing the aggregate dynamics of a TFP

¹See for example Violante, Kaplan, and Weidner (2014), Auclert (2017) and Manovskii and Hryshko (2017). Recent work by Arellano, Blundell, and Bonhomme (2017) uses a quantile-based approach to allow for non-linear dynamics, but the time aggregation problem is not addressed.

²While access is very restricted, panel data from financial aggregation platforms has been highly informative about consumption behavior. In the US examples include Gelman, Kariv, Shapiro, Silverman, and Tadelis (2014), Ganong and Noel (2017) and Baker (2015), while Vardardottir and Pagel (2016) uses data from Iceland.

shock were not much changed in a heterogeneous agent model, have resulted in a slow start for this literature. However, recent advances have allowed for a new generation of Heterogeneous Agent New Keynesian (HANK) models that, as their name suggests, combine elements from both the heterogeneous agent and New Keynesian literature. These models not only match the growing evidence on micro consumption behavior, but also imply very different aggregate dynamics and/or propagation mechanisms following macroeconomic shocks, compared to their representative agent equivalents. In particular the transmission mechanism of monetary policy can look very different in a HANK model. For example, in the model of Kaplan, Moll, and Violante (2016) the intertemporal substitution channel is dwarfed by indirect general equilibrium effects, in stark contrast to a representative agent model.

While these HANK models make clear the potential importance of heterogeneity in economic fluctuations, and particularly for monetary policy, their quantitative results hinge on previously unimportant assumptions such as the tenure of debt instruments and the government’s fiscal rule. Thus far the ability of these models to help us distinguish transmission channels empirically has been limited. Auclert (2017), in contrast to the fully structural HANK models, takes a simplified approach to aggregate dynamics, and one that we will follow in this paper, in order to derive a set of sufficient statistics, directly measurable from a suitable dataset, that is highly informative about the relative size of different monetary policy transmission channels. His methodology benefits from being transparent and closely tied to the data. However, with the small sample sizes of the publically available panel surveys he uses, along with the methods used to impute transitory MPCs (including BPP), he is in the end unable to attain convincing quantitative estimates for the sufficient statistics he identifies.³ Our new methodology for estimating consumption behavior, combined with the fact that our panel covers the entire Danish population so that we can reconcile aggregate balance sheet positions of households with that of the government, firms and foreigners in the national accounts, provides an almost ideal setting in which to estimate Auclert’s sufficient statistics. We find heterogeneity in households’ interest rate exposure to be very relevant to monetary policy transmission. Households that stand to lose from an interest rate hike (typically those with a large variable rate mortgage) are likely to cut back their consumption far more than those who stand to gain (typically those with large liquid asset holdings). We estimate this interest rate exposure channel to be of an order of magnitude larger than the intertemporal substitution channel.

1.1 Empirical Evidence on Heterogeneity in Consumption Behavior

Most micro empirical evidence on consumption behavior comes in the form of an estimate of the marginal propensity to consume from a one-off source of income over the following three months to one year. Table 1 shows a selection of the population average estimates from the literature. In representative agent models of the macroeconomy, the agent is usually found to have an annual MPC out of transitory shocks of around 3-5%, about an

³Fagereng, Holm, and Natvik (2016) also estimate these sufficient statistics, imputing MPCs from lottery winnings in Norway, but they are also limited by sample size.

order or magnitude lower than most of the empirical estimates in this table. The inability of these macro models to match even this most basic fact about micro consumption behavior undermines their claim to be micro-founded and raises questions about the reliability of their macro implications. Ad-hoc fixes, such as the addition of the hand-to-mouth consumers in Campbell and Mankiw (1989), also bear little relation to the micro data. The new set of heterogeneous agent models come with testable implications for the distribution of MPCs across dimensions such as net or liquid wealth (see Carroll, Slacalek, Tokuoka, and White (2016) and Violante, Kaplan, and Weidner (2014)). Most of the studies in table 1 do not have the power to say much if anything along these dimensions. Three methods are used to empirically determine the marginal propensity to consume. The first is to identify a natural experiment and then measure the consumption response to it. Often this is done using the Consumer Expenditure Survey in the US. For example Johnson, Parker, and Souleles (2006) use randomly assigned timing of 2001 tax rebates and specially included questions in the Consumer Expenditure Survey to identify a three month aggregate marginal propensity to consume of 0.2-0.4. This method is generally considered the most reliable, but estimates vary and there is no strong consensus. Identification issues arise as to when exactly households learn about the payment versus when it is received and it is unclear the extent to which external validity extends from these natural experiments to the kinds of transitory shocks found in heterogeneous agent models. As most of these studies rely on consumer survey data they tend to lack power due to high measurement error and low sample sizes. As a result they have produced very little evidence of how the MPC varies among different groups in the economy. A recent paper by Fagereng, Holm, and Natvik (2016) overcomes some of these problems. By using lottery data, the shock to income is truly random.⁴ They use registry data from Norway similar to the data we use from Denmark and have a sample of over 30,000 lottery winners over 10 years. As a result they can identify the MPC for households with differing liquid wealth, as well as by the size of the lottery win. They find that households in the lowest quartile of liquid wealth have an MPC of approximately 0.61 over a 6 month period, as opposed to 0.45 for households in the highest quartile of liquid wealth. Using data from a financial aggregator, Gelman (2016) has enough power to identify large differences in the impulse response to a tax rebate at a monthly frequency between household quintiles of cash-on-hand.

The second method is simply to ask individuals how much of a transitory income change they would consume, as was done in the Italian Survey of Household Income and Wealth in 2010 and the NY Fed’s Survey of Consumer Expectations in 2016-2017. Jappelli and Pistaferri (2014) find an aggregate MPC of 0.48 using this Italian data and are able to identify clear differences across levels of liquid wealth. Fuster, Kaplan, and Zafar (2018) find a lower aggregate MPC in the NY Fed’s survey, but find heterogeneity by both size and sign of the shock. While this method holds great promise, it is clearly limited by the accuracy of households’ own response to the question.

⁴We should note that even lottery winnings have some problems. First the results hold for winners of the lottery who may not be representative of the wider population. Second the consumption response to a lottery win may be very different to other income shocks. For example you may spend a significant portion of a small lottery win just celebrating the fact.

Table 1 Estimates of the Marginal Propensity to Consume from Income Shocks

Permanent Shocks	Consumption Measure				Method	Event /Sample
	Nondurables	Total PCE	Horizon			
Blundell, Pistaferri, and Preston (2008)* Gelman, Gorodnichenko, Kariv, Koustas, Shapiro, Silverman, and Tadelis (2016) Transitory Shocks	0.65	1.0	~ ~	1 3	1 3	Estimation Sample: 1980–92 Gasoline Price Shock
Agarwal and Qian (2014)		0.90	10m	1		Growth Dividend Program Singapore 2011
Blundell, Pistaferri, and Preston (2008)*	0.05	~ 0		3		Estimation Sample: 1980–92
Browning and Collado (2001)				1		Spanish ECPF Data, 1985–95
Coronado, Lupton, and Sheiner (2005)		0.36	1y	1		2003 Tax Cut
Fuster, Kaplan, and Zafar (2018)		0.08–0.31	3m	2		NY Fed Survey Cons. Expectations
Gelman (2016)		0.13	3m	1		Tax refunds 2013–2016
Hausman (2012)		0.6–0.75	1y	1		1936 Veterans' Bonus
Hsieh (2003)*	~ 0	0.6–0.75		1		CEX, 1980–2001
Jappelli and Pistaferri (2014)	0.48			2		Italy, 2010
Johnson, Parker, and Souleles (2009)	~ 0.25		3m	1		2003 Child Tax Credit
Lusardi (1996)*	0.2–0.5			3		Estimation Sample: 1980–87
Parker (1999)	0.2		3m	1		Estimation Sample: 1980–93
Parker, Souleles, Johnson, and McClelland (2013)	0.12–0.30	0.50–0.90	3m	1		2008 Economic Stimulus
Sahm, Shapiro, and Slemrod (2010)		~ 1/3	1y	1		2008 Economic Stimulus
Shapiro and Slemrod (2009)		~ 1/3	1y	1		2008 Economic Stimulus
Souleles (1999)	0.045–0.09	0.34–0.64	3m	1		Estimation Sample: 1980–91
Souleles (2002)	0.6–0.9		1y	1		The Reagan Tax Cuts of the Early 1980s

* Elasticity. Methods: 1) Natural Experiment 2) Survey question 3) Covariance restrictions

The third method is to impose covariance restrictions on panel data of income and consumption and use these to identify the consumption response to income shocks of differing persistence. The most well known of these is the paper by Blundell, Pistaferri, and Preston (2008) which uses imputed non-durable consumption data based on food expenditure reported in PSID data. They estimate a consumption elasticity (closely related to an MPC if households consume most of their income) and find very little consumption response to transitory shocks, but as discussed previously this estimate is strongly downward biased.

This paper also adds to the limited literature on consumption responses to permanent shocks to income. Natural experiments for permanent shocks are very hard to come by. Gelman, Gorodnichenko, Kariv, Koustas, Shapiro, Silverman, and Tadelis (2016) use shocks to gasoline prices as a proxy for a permanent shock to income and find an MPC close to 1 across the population. BPP find a consumption elasticity to permanent shocks to income around 0.65 (the permanent shock elasticity is less affected by the time aggregation problem). For a more complete overview of the literature on consumption responses to income changes, see Jappelli and Pistaferri (2010).

2 Empirical Strategy

2.1 Methodology Intuition

Our empirical strategy will be to impose some covariance restrictions on our data and apply GMM in order to estimate parameters of interest. While this strategy allows us to precisely estimate these quantities, in some ways it obscures from the key features of the data that are driving the result. In this section we present some very simple regressions of expenditure growth on income growth and compare them with what we would expect in some very well understood baseline models.

We will look at the estimate of β^N in the model

$$\Delta^N c_{it} = \alpha^N + \beta \Delta^N y_{it} + \varepsilon_{it}$$

where N , the number of years over which growth is measured, varies from 1 to 10. Our identification comes from the fact that transitory income shocks make up a relatively large proportion of the variance of income growth over a short period, and that permanent income shocks dominate the variance of income growth over a long period. Figure 1 shows what we would expect to see under three well known models, as well as what we actually observe the data. In a complete markets model in which all idiosyncratic shocks to income are insured against there is no relation between income growth and consumption growth, as represented by the blue horizontal line at zero. In the Solow growth model, and also some old Keynesian models, households' expenditure is a constant proportion of income that period, regardless of transitory shocks to income. The green horizontal line around 0.75 shows what we would see in a model of this type where households spend 75% of their income each period. The red line shows the results for a typical buffer-stock saving model. In this model the regression of consumption




Figure 1 Regression coefficients of consumption growth on income growth

growth on income growth over one year yields a relatively small number as households are able to self-insure against the transitory shocks that dominate at this frequency. As the time period over which income growth is measured increases the observed income growth is proportionally more permanent and self-insurance is not possible. The red line asymptotes towards 1.0 as N gets large.

The gray line shows the results of these regressions using all households in the Danish sample. It is striking that the data appears to be closest to the Solow model, with only a small decrease in the regression coefficient over short periods. However, aggregating all households in this way hides a large degree of heterogeneity, particularly across households with different levels of liquid wealth. The two black lines shows the regression coefficients where the sample is restricted to households in the lowest and highest quintiles of liquid wealth (averaged over the observed period) respectively. For households in the lowest quintile there is no evidence of consumption smoothing. As the regression coefficient is relatively stable for this group over N , the result that the marginal propensity to consume out of both transitory and permanent shocks are similar and very high is robust to a large degree of model misspecification in the next section. Households in the top quintile of liquid wealth show a clear upward slope in figure 1, indicating a

substantial degree of consumption smoothing. The fact that the regression coefficient for this group appears to asymptote well below 1 also suggests, in contrast to standard buffer-stock models, that the MPC out of permanent shocks for liquid households is significantly lower than 1.

2.2 Aside: Why Not BPP? A Brief Introduction to the Time Aggregation Problem

As explained above, our identification is going to come from the shape of income and consumption covariance over increasing periods of time. An obvious question is why we have chosen not to use the well known methodology of Blundell, Pistaferri, and Preston (2008) who achieve identification of transitory shocks from the fact that a transitory shock in period t will mean-revert in period $t + 1$.⁵ Unfortunately the method is not robust to the time aggregation problem of Working (1960). While macro-economists have long been aware of the importance of time aggregation in time series regressions (see Campbell and Mankiw (1989) for a well known example), the problem appears to have escaped the attention of the household finance and labor economics literature. We will therefore briefly describe the problem here. For a more detailed account with particular attention to BPP, see Crawley (2018).

Figure 2 shows how the problem arises. The solid ‘Income flow’ line shows the true income flow of a household who receives zero income throughout year 1, zero income for the first half of year two, and then a constant income flow of 1.0 per year during the second half of the second year and in the third year. The dashed line shows the observed *total* income of the household in years 1, 2 and 3: zero in year 1, 0.5 in year 2 and 1.0 in year 3. The important thing to note is that despite there only being one ‘shock’ to the income flow over the three year period, the naïve observer would assume there had been two shocks, one between years 1 and 2 and another between years 2 and 3. This effect is of particular importance to econometric techniques that make use of the auto-covariance structure of data processes. For example the first difference of a random walk in discrete time has no autocorrelation, but the first difference of a time-aggregated random walk in continuous time has an autocorrelation equal to $1/4$. BPP uses time aggregated income data and achieves identification of transitory variance precisely through the auto-covariance structure. This is why, as explained in the introduction, the problem is particularly pervasive for this methodology.

While it would be possible to stick very closely to the original BPP model and adjust the covariance restrictions to take account of the time aggregation problem,⁶ we have found that in practice the underlying assumptions made by BPP (in particular that

⁵Kaplan and Violante (2010) show in discrete time simulations that the methodology works reasonably well for standard calibrations of buffer-stock models and end up concluding “The BPP insurance coefficients should become central in quantitative macroeconomics”. However, some recent papers such as Commault (2017) and Hryshko and Manovskii (2018) have pointed to other potential problems of the methodology.

⁶Crawley (2018) takes this more straightforward approach using the same PSID data as used in BPP.

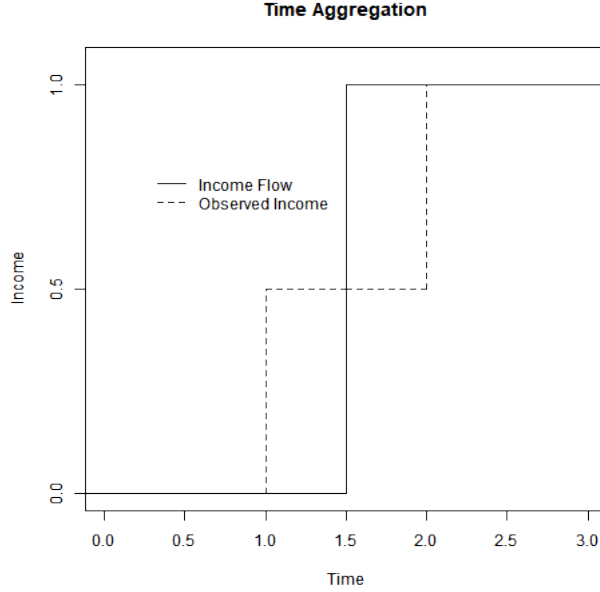


Figure 2 The Time Aggregation Problem

consumption follows a random walk) do not fit with the data.⁷ Therefore we have chosen to attain identification in a manner similar to Carroll and Samwick (1997) which allows us to be agnostic about the exact short term dynamics of income and consumption.

2.3 Covariance Restrictions

2.3.1 Income Dynamics: Carroll and Samwick (1997)

Our identification of permanent and transitory income variance will follow the methodology of Carroll and Samwick (1997) closely. Our method will correctly account for time aggregation, but due to identification coming from income growth over 3, 4 and 5 years, rather than the covariance of income growth at time $t + 1$ with time t as in BPP, time aggregation only introduces a small bias in the estimates of Carroll and Samwick (1997). We will first describe the method without time aggregation and then show how the estimates need to be adjusted.

Carroll and Samwick (1997) assume that income is composed of a permanent component that follows a random walk and a transitory MA(2) component. That is:

$$\begin{aligned} y_t &= p_t + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} \\ p_t &= p_{t-1} + \zeta_t \end{aligned}$$

⁷Kaplan and Violante (2010) show that without time aggregation, the BPP method correctly identifies the transitory consumption response in the period of the income shock regardless of the consumption dynamics going forward. This fact is again not robust to the time aggregation problem. With time aggregation taken into account the estimates are highly sensitive to assumptions about short term consumption dynamics.

where ζ_t and ε_t are mean zero random variables, independent of each other and of themselves over time. Each has constant variance, σ_ζ and σ_ε respectively. For $N \geq 3$ we have:

$$\begin{aligned} \Delta^N y_t &= \zeta_t + \zeta_{t-1} + \dots + \zeta_{t-N+1} \\ &\quad + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} - (\varepsilon_{t-N} + \theta_1 \varepsilon_{t-1-N} + \theta_2 \varepsilon_{t-2-N}) \\ \Rightarrow \text{Var}(\Delta^N y_t) &= N\sigma_\zeta + 2 \underbrace{(1 + \theta_1^2 + \theta_2^2)\sigma_\varepsilon}_{\text{"Total" transitory variance}} \quad \text{for } N \geq 3 \end{aligned} \quad (1)$$

Equation 1 shows that the variance of income growth grows linearly with the number of years of growth beyond 3. The transitory component adds variance at the beginning and end of the growth period, but any transitory shock to income that occurs in the middle of the period does not affect income growth as it will have died out by the end. Carroll and Samwick (1997) use this to identify the variance of permanent shocks ($\text{Var}(\zeta)$) and the “total” transitory variance $((1 + \theta_1^2 + \theta_2^2)\text{Var}(\varepsilon))$. While similar to BPP it is important to note that BPP attempts to identify the variance of initial impact of the transitory shock, $\text{Var}(\varepsilon)$, rather than the “total” transitory variance. While the notion of “total” transitory variance will carry over naturally into the time aggregated case, the variance of the initial impact does not have a natural interpretation.

The administrative data we use in this paper is at an annual frequency and measures the sum of income over the observed year. If shocks to income occurred only on 1st January every year then we could use equation 1 to identify permanent and transitory variance. It is important to distinguish between a model in which shocks happen about once a year (for example) but can occur at any point in the year, versus a model in which shocks to income happen on a timetable. The former can be modeled in continuous time with jumps occurring as a Poisson process approximately once a year. The later is best modeled as a discrete time model. In this paper we will take the former approach. While some types of jobs may have a regular schedule on which pay appraisals take place, many of the larger permanent shocks to income occur when a worker changes job which can occur at any point in the year. We (along with the literature) lack a clear understanding of what makes up the bulk of the transitory shocks to income and the Danish data is a potentially rich source for further research in this area.⁸ Furthermore, even if each individual household experienced shocks on a pre-set timetable, if the timetable itself varies across the year for different households, our approach would yield unbiased results. While there is a big change in going from an underlying annual process to a quarterly process, the further change from quarterly to continuous time is much smaller.⁹ As the exposition is much simpler in continuous time, we will therefore chose to present our own method in continuous time.

To write the equivalent model in continuous time we define two underlying martingale processes (possibly with jumps), P_t and Q_t where P_t will represent the *flow* of permanent income at time t and the change in Q_t provides the transitory *impulses* that generate

⁸While we use annual income data in this paper in order to match with our expenditure data, monthly labor income data is collected, along with employer-employee matching.

⁹See Crawley (2018)




Figure 3 Generic Transitory Shock Impulse Response

the transitory income. We assume that for all $s_1 > s_2 > s_3 > s_4 > 0$:

$$\begin{aligned}\text{Var}(P_{s_1} - P_{s_2}) &= (s_1 - s_2)\sigma_P^2 \\ \text{Cov}(P_{s_1} - P_{s_2}, P_{s_3} - P_{s_4}) &= 0 \\ P_s &= 0 \quad \text{if } s < 0\end{aligned}$$

and similarly for Q_t . That is these martingales have independent increments. As a useful benchmark, two Brownian motions satisfy these criteria.

The natural generalization of the MA(2) transitory income process from Carroll and Samwick (1997) is to allow for a generically shaped transitory income shock that decays to zero in under 2 years. Figure 3 shows an example of such a transitory income shape $f(t)$, but the model also allows for completely transitory shocks in which case $f(t)$ would be a delta function with all the income from the transitory shocks arriving as a mass at the time of the shock. In this model the *flow* of income arriving at time t is given by the flow of permanent income and the sum of income arising from any transitory shocks to income that have occurred in the previous two years:

$$y_t = P_t + \int_{t-2}^t f(t-s)dQ_s$$

We do not observe y_t directly but instead \bar{y}_T , the time aggregated income over each one year period.

$$\bar{y}_T = \int_{T-1}^T y_t dt \text{ for } T \in \{1, 2, 3, \dots\} \quad (2)$$

Taking the N^{th} difference for $N \geq 3$ we get:

$$\Delta^N \bar{y}_T = \int_{T-1}^T y_t dt - \int_{T-N-1}^{T-N} y_t dt$$

$$\begin{aligned}
&= \int_{T-1}^T (T-s)dP_s + (P_{T-1} - P_{T-N}) + \int_{T-N-1}^{T-N} (s - (T-2))dP_s \\
&\quad + \left(\int_{T-1}^T \int_{t-2}^t f(t-s)dQ_t - \int_{T-N-1}^{T-N} \int_{t-2}^t f(t-s)dQ_t \right)
\end{aligned} \tag{3}$$

The variance of time aggregated income of an N year period is therefore:¹⁰

$$\text{Var}(\Delta^N \bar{y}_T) = (N - \frac{1}{3})\sigma_P^2 + 2\text{Var}(\tilde{y}) \text{ for } n \geq 3 \tag{4}$$

This is similar to the non time aggregated case (equation 1) except the coefficient on permanent variance is $N - \frac{1}{3}$. This error, though having less serious consequences than for BPP, has nevertheless been overlooked by the large literature that studies income dynamics using panel data.¹¹ As with the MA(2) case the transitory variance identified is the variance of “total” transitory income received in the year, \tilde{y} , where this is defined as:

$$\tilde{y}_T = \int_{T-1}^T \int_{t-2}^t f(t-s)dQ_s dt \tag{5}$$

We will use equation 4 with $N \in \{3, 4, 5\}$.

2.3.2 Consumption Dynamics

Our approach will be to extend the identification of income variance by using growth over 3, 4 and 5 years to also identify the covariance of income and consumption. In contrast to BPP, which assumes consumption follows a random walk, we will instead assume that the impulse response to a transitory shock follows a generic path, $g(t)$, that like the transitory income shock has fallen to zero two years after the news of the shock. Figure 4 shows possible paths for both income and consumption, along with the alternative random walk impulse response of BPP. The best evidence for the speed at which the consumption response decays comes from Gelman (2016) and Fagereng, Holm, and Natvik (2016), both of which show the response has entirely or almost entirely decayed two years after the shock. In section 6.3 we will show how this assumption may potentially biases the transitory consumption response up but that this bias is small, especially for all but the most liquid households. We will maintain the assumption from BPP that the consumption response to a permanent shock to income follows a random walk proportional to the permanent shock. Under these assumptions the instantaneous flow of consumption is given by:

$$c_t = \phi P_s + \int_{t-2}^t g(t-s)dQ_s$$

¹⁰See appendix A for full details of this derivation, including how we can approximate a log income process with levels.

¹¹For examples see Moffitt and Gottschalk (2012), Meghir and Pistaferri (2004), Nielsen and Vissing-jorgensen (2004), Heathcote, Perri, and Violante (2010) and more recent quantile regression approaches such as Arellano, Blundell, and Bonhomme (2017).




Figure 4 Generic Transitory Shock Impulse Response

and the covariance of time aggregated income and consumption growth over $N \geq 3$ years is given by

$$\text{Cov}(\Delta^N \bar{c}_T, \Delta^N \bar{y}_T) = \phi(N - \frac{1}{3})\sigma_p^2 + 2\text{Cov}(\tilde{c}, \tilde{y}) \text{ for } N \geq 3 \quad (6)$$

where total transitory income, \tilde{y} is given by equation 5 and total transitory consumption, \tilde{c} , is defined by:

$$\tilde{c}_T = \int_{T-1}^T \int_{t-2}^t g(t-s) dQ_s dt \quad (7)$$

Using the equations for variance (4) and covariance (6) of income and consumption growth over N years for at least 2 different values of N , we are able to identify the 4 unknowns we are interested in:

1. σ_p^2 Variance of permanent shocks
2. $\sigma_q^2 = \text{Var}(\tilde{y})$ Variance of transitory income received in a year
3. ϕ Marginal Propensity to eXpend (MPX) w.r.t permanent income
4. $\psi = \frac{\text{Cov}(\tilde{c}, \tilde{y})}{\text{Var}(\tilde{y})}$ Regression coefficient of transitory consumption w.r.t transitory income over a year (MPX out of transitory income)

Our panel data covers 13 years and we choose to use growth over 3, 4 and 5 years to balance greater identification (longer growth periods give more power) with three identification problems that grow with N . The first is the fact that many households drop out of the sample if we demand they have reliable data for the too many consecutive years. The second is that if the permanent shock in fact decays slowly over time (e.g. is in fact AR(1)), the bias this introduces will be larger for large N . Finally, the validity of running the regressions in levels (rather than logs) is reduced over large N when the potential for the variance of income to change significantly from the start to end of the

sample is high. In section 6.3 we test the importance of these issues. Using growth over 3 different time periods means the system is over identified with 6 equations with which to estimate these 4 unknowns. We follow BPP and use diagonally weighted minimum distance estimation, although our results are not significantly changed by using other popular weighting methods.¹²

As a lot of our analysis will focus on the parameter ψ is it worth describing exactly what this is and why we have labeled it the marginal propensity to expend out of transitory income. If we were able to exactly observe transitory income and consumption resulting from transitory income then ψ would be the regression coefficient of on this transitory consumption on transitory income. If transitory income shocks have no persistence this is approximately a six month MPX (on average the shock will happen six months into the year so that the regression will pick up the change in consumption in the following six months. If transitory income shocks have a little persistence (appendix B shows evidence of a small amount of transitory income persistence) ψ can only loosely be interpreted as the MPX to an income shock and the reader should bear in mind that the true interpretation is, "if income is higher by one unit this year due to transitory factors, then consumption this year will be expected to be higher by ψ units".

3 Data

Our panel data on income and expenditure comes from Danish registry data from 2003-2015. This data has a number of advantages over survey based measures. First, the sample contains millions of households rather than thousands. Second, households are required by law to report their data so there is much less risk of selection bias through drop outs. Third, measurement error in income data is largely eradicated, as employees' income data is third party reported by their employer, compared to survey data where self reported income has been shown to be particularly unreliable for irregular income.¹³

3.1 Income

We are interested in income and consumption decisions at the household level. We define a household as having either one or two adult members. Two adults are considered to be in the same household if they are living together and a) are married to each other or have entered into a registered partnership, b) have at least one common child registered in the Civil Registration System or, c) are of opposite sex and have an age difference of 15 years or less, are not closely related and live in a household with no other adults.¹⁴

¹²In general our results may be subject to misspecification problems, but the sample size of our data means that standard errors are small.

¹³See David, Marquis, Moore, Stinson, and Welniak (1997) for a survey of income measurement error issues in survey data.

¹⁴Adults living at the same address but not meeting one of the three criteria are regarded as separate families. Children living with their parents are regarded as members of their parents' family if they are under 25 years old, have never been married or entered into a registered partnership and do not themselves have children. A family meeting these criteria can consist of only two generations. If three or more generations live at the same address, the two younger generations are considered one family, while the members of the eldest generation constitute a separate family.

In the panel data, an individual's household will change if he or she gets married or divorced. This leads to some selection bias given that we require households to survive for at least 5 years. Following the literature our baseline results will be reported using the labor income of the head of household.¹⁵ We will use after tax and transfer income as we are interested in the consumption response to these changes in income, although the method could be used to measure the extent of consumption insurance provided by the tax and transfer system. Our data comes from the administrative records from the tax authority. The tax reporting system in Denmark is highly automated and individuals bear little of the reporting burden. For employees income is reported by their employers and is thought to be highly accurate. The underground economy in Denmark is small. We remove business owners from the sample as their income may be less accurately reported, but more importantly, because the expenditure imputation method does not work well for them (see section 3.2).

We work with the residual of income after controlling for observable characteristics of households that may affect their income and consumption. To start we remove households in the top and bottom 1% of the income distribution. We then normalize by average household income over the observed period, and regress income on dummies for age, year, highest level of education, marital status, homeowner status and number of children along with interaction of age with education, marital status and homeowner status. We take the change in the residuals of this regression to be the unexpected income change for a household from one year to the next and remove households in the top and bottom 1% of the unexpected income *change* distribution.

3.2 Imputed Expenditure

Our expenditure data comes from imputing expenditure from income and wealth. Along with other Scandinavian countries, Denmark is unusual in that tax reporting includes information about wealth along with income, a legacy from the wealth tax that was phased out between 1989 and 1997. Following the methodology from Browning and Leth-Petersen (2003) and Fagereng and Halvorsen (2015) we impute expenditure using the identity:

$$\bar{C}_t \equiv \bar{Y}_t - \bar{S}_t = \bar{Y}_t - P_t - \Delta NW_t$$

where \bar{C}_t , \bar{Y}_t and \bar{S}_t are the sum of expenditure, income and savings over the year t respectively. P_t is contributions to privately administered pension schemes, for which we have very accurate data due to tax deductibility ΔNW_t is the change in (non-pension, non-housing) net worth measured at the end of years t and $t - 1$. Banks and brokers are required to report the value of their clients' accounts on 31st December each year, and the tax reporting year runs from 1st January to 31st December, so the data for income and wealth reported in the tax returns matches with that required to use this identity to impute consumption.

¹⁵See Moffitt and Zhang (2018) for an overview of the literature on income volatility in the PSID. In contrast the PSID literature we define the head of household as the highest earner over the 13 year period in our sample. We believe this definition better fits the social structure in Denmark.

The method works well for households with simple financial lives. One of the biggest problems with the method is its inability to handle capital gains well. The income used in the imputation includes all labor income and capital income, however it excludes capital gains. The value of assets will vary both due to savings from reported income but also due to capital gains and losses. We handle this in a number of ways. First, we completely exclude housing wealth from our measures of net worth and saving, treating housing as an off balance sheet asset. The problem with treating housing in this way is that we must exclude households in years in which they are involved in a housing transaction. For the self employed, it is also difficult to distinguish between expenditure and investment in their business, so we exclude all households who receive more than a trivial amount of their income from business ventures. Finally, households that hold significant equity investments are likely to see sizable capital gains and losses. We make a naive adjustment by making the assumption that they hold a diversified index of stocks. While this will likely lead to significant measurement error for these individuals, the concern is mitigated first by the fact that stock holding is much more unusual in Denmark than in the US for example. Only around 10% of households hold any stocks, and for many of those stocks make up only a small proportion of their total wealth. Furthermore, as we will explain in section 6.4, measurement error in consumption is not a concern unless it is correlated with income. This seems unlikely to be the case, except for households that hold significant equity in the firms in which they work. Another concern with the imputation method is transfers of wealth, say between family members or friends. Indeed imputed expenditure is negative for approximately 3% of households and this may explain a proportion of that. We throw out both income and expenditure data for households in years in which their expenditure is negative. In section 6.3 we test the robustness of our results to sample selection bias problems that these issues may give rise to.

As with income, we work with the residual of expenditure after normalizing by mean household income and controlling for the same observable features as income. We follow exactly the same steps as described in section 3.1.

In evaluating how much we can learn from such a measure, it should be compared to the best alternatives available to economists. In the original BPP paper the authors only have access to food expenditures from the PSID data and impute total non-durable consumption by comparison with the Consumer Expenditure Survey. Self reported consumption is also notoriously poor quality even in comparison to self reported income. Furthermore, in the PSID data the questions in food expenditure are ambiguous as to which period exactly the question is referring to. A recent paper by Abildgren, Kuchler, Solange, and Sorensen (2018) shows that the mean levels of expenditure from this imputation method are close to those from the national accounts (see figure 5). They find large differences at the household level between the consumer survey and imputed expenditure although it is not clear that this is a problem with the imputation method as opposed to the survey measure. Indeed for car purchases, for which we have highly accurate register data, the consumer survey shows significant underreporting, consistent with Koijen, Nieuwerburgh, and Vestman (2014) who find 30% underreporting of car purchases in the Swedish consumer survey. We believe that, with the exception of

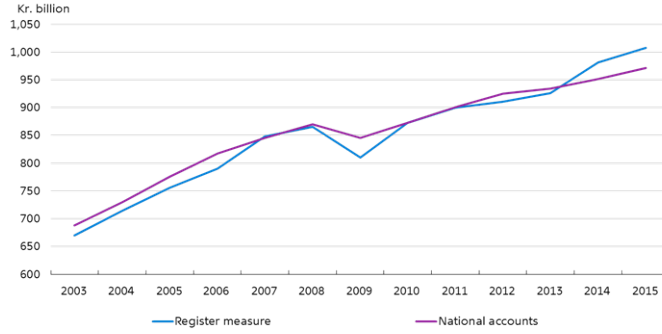


Figure 5 Imputed Register Measure and National Account Measure of Expenditure
(from Abildgren, Kuchler, Solange, and Sorensen (2018))

transaction level data reported by financial aggregation applications, the imputation method we use is some of the highest quality expenditure data available to researchers for the types of questions we are addressing.

3.3 Sample Selection

As our methodology requires income uncertainty to be relatively constant through the observed period¹⁶ and the young and old are likely to have significant predictable income trends unobservable to the econometrician, we limit the sample to households headed by an individual between the age of 30 and 55 in 2008. Our final sample contains 7.7 million observations over 2004-2015 from an age group population total of 18.1 million. The selection criteria that reduces the sample size the most is the requirement that a household does not make a housing transaction for a period of 5 years. Table 3.3 shows summary statistics for all Danish households whose head fits into this age group as a whole as well as the sample we use in estimation. It is reassuring that both the mean and median values for after tax income and consumption are similar in the estimation sample and the population. Our estimation sample has much lower standard deviations as a mechanical result of excluding the top and bottom 1% of the income and consumption distributions which contains extreme values. Sample selection shows up in homeownership and car ownership as we exclude those households buy a house at the end of a five year period but who otherwise would be counted as renters. This also results in our sample being on average one year older than the population. Unhedged Interest Rate Exposure (URE) and Net Nominal Position (NNP) will be discussed in section 5, but again the significant differences here are due to the housing transaction criteria.

¹⁶Appendix B shows the assumption holds for this age group.

	Estimation Sample			Population (Age 30-55)		
	Mean	Median	Std Dev	Mean	Median	Std Dev
After Tax Income	59,261	57,804	28,819	58,312	53,304	68,799
Consumption	52,680	48,344	28,581	51,427	45,222	43,312
Liquid Assets	18,438	6,856	33,016	23,331	6,578	81,473
Net Worth	74,937	19,115	157,295	85,799	12,952	564,404
Homeowner	0.57	1.00	0.50	0.50	1.00	0.50
Car Owner	0.66	1.00	0.47	0.55	1.00	0.50
Higher Education	0.31	0.00	0.46	0.33	0.00	0.47
Age	43.5	44.0	7.1	42.5	42.0	7.3
URE	-28,052	-12,627	108,382	-44,134	-18,395	245,655
NNP	-109,685	-65,810	156,523	-158,321	-85,207	542,498
No. household-year obs	7,664,360			18,050,340		

Notes: Values are 2015 USD. Age refers to the age in 2008 of the main income earner in the household. For the purposes of calculation of consumption in the population, top and bottom 1% in terms of consumption have been excluded. URE and NNP can only be calculated in the period 2009-2015 due to mortgage information being insufficiently detailed in the previous years.

Table 2 Summary Statistics

4 Income and Consumption Characteristics by Household Wealth

Using our entire estimation sample we find a mean MPX out of transitory shocks of 0.50 and a mean MPX out of permanent shocks of 0.72. However, these averaged results hide a significant amount of heterogeneity. From the standpoint of consumption theory it is the ability of households to self-insure with their own wealth that most determines how much they smooth their consumption over shocks. We divide our estimation sample into quintiles according to both liquid wealth (which we define as bank deposits¹⁷) and net wealth. In each case wealth is measured as the mean household wealth holdings over the entire sample period.

Figure 6 shows the estimated income variances and MPX's for households in each quintile of liquid wealth. Looking at the left hand variance panel first, it is noticeable that income uncertainty, and particularly permanent income uncertainty, is highest for households in the lowest quintile of liquid wealth. This is some evidence towards the idea that heterogeneous tastes (e.g. discount factors of risk aversion) may be more important than income risk in determining wealth held for precautionary saving. For households in the top three quintiles of liquid wealth it is remarkable how similar their level of income risk is. Note that in contrast to standard estimates of the US income process, permanent income variance in Denmark is slightly higher than transitory variance, likely due to the high levels of social insurance available in Denmark. The variance level, at just over

¹⁷ The results are little changed using any other definition of liquid wealth as long as housing and debts are excluded. See section 6.3.




Figure 6 Variance and MPX by Liquid Wealth Quintile

0.002 for these top three quintiles, represents a standard deviation of just below 5% of permanent income per year.

Note that the estimates of income variance we obtain are highly sensitive to our treatment of outliers, but our MPX estimates do not change.¹⁸

The right hand panel of figure 6 shows our estimates for the MPX out of permanent and transitory shocks by liquid wealth quintile. The lowest wealth quintile, who hold less than \$2,000 in bank deposits on average over the sample period, look somewhat like hand-to-mouth consumers. They respond almost equally to permanent and transitory shocks, spending over 80% of income shocks in the year that it arrives. However, the fact that both permanent and transitory MPX's are very similar and significantly less than 1 suggests these households may be more accurately modeled as saving in an illiquid asset such as housing or a pension following a rule of thumb (say 20% of income) and then living hand to mouth on the remainder. As the quintile of liquid wealth increases, the MPX out of both transitory and permanent income decreases. In the top quintile, formed of households that maintained a mean bank balance above \$30,000, the MPX out of permanent shocks is 0.57 and out of transitory shocks 0.23. From the point of view of theory the responsiveness of spending out of permanent shocks in this quintile is low, while that of transitory shocks is high. A more thorough discussion of how these results compare to a standard model calibrated to Danish characteristics will wait until section 6.

Figure 7 shows the estimates for households grouped by quintiles of net wealth. Here the pattern is slightly different. The quintile with the highest MPX out of both transitory and permanent income is the second lowest, the quintile that contains zero net worth. Households in the lowest quintile, those with over \$20,000 in net debt, do not seem

¹⁸See section 6.3 for evidence of this as well as a discussion of why this is the case.

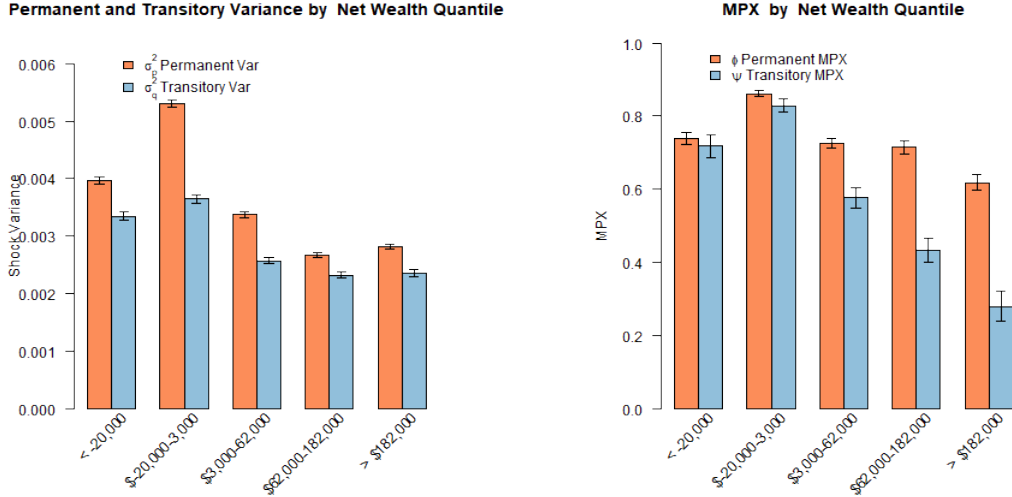


Figure 7 Variance and MPX by Net Wealth Quintile

to distinguish between permanent and transitory income shocks in their consumption responses, but their MPX for both is about 10 percentage points lower than the quintile with close to zero net wealth. The pattern for quintiles 3 to 5 looks similar to that for liquid wealth: the MPX out of transitory shocks falls sharply to around 0.28, while that out of permanent shocks also falls but more slowly to 0.62.

These results are broadly in line with the literature. The population mean of 0.5 for transitory MPX is a little higher than most estimates from table 1, but bearing in mind our estimate includes durables and is best compared to a six month MPC, it is certainly not an outlier. The MPX out of permanent shocks of 0.72 is also between the BPP estimate of 0.65¹⁹ and the estimate of 1.0 from Gelman, Gorodnichenko, Kariv, Koustas, Shapiro, Silverman, and Tadelis (2016). The strength of the relationship between liquid wealth and MPC is similar to that found in Gelman (2016) and stronger than in Fagereng, Holm, and Natvik (2016).

5 Application: Monetary Policy and the Redistribution Channel

Auclert (2017) lays out a clear and intuitive theory as to how heterogeneity in the marginal propensity to consume out of transitory shocks affects the transmission mechanism of monetary policy. He identifies five channels through which monetary policy can act, three of which are mute in the absence of heterogeneity. He then uses this theory to identify a small set of sufficient statistics that help distinguish which of these

¹⁹The permanent ‘insurance’ coefficient estimated by BPP does not suffer as much from the time aggregation problem as the transitory coefficient.

channels are of quantitative importance. While these statistics in theory are highly informative about the transmission mechanism of monetary policy, in his paper he has neither the data nor the methods to be able to estimate them convincingly. He states, “As administrative quality household surveys become available and more sophisticated identification methods for MPCs arise, a priority for future work is to refine the estimates I provide here”. Given we have administrative data, along with a new method to estimate MPCs, a natural application of our work is to estimate Auclert’s sufficient statistics. Our data has two significant advantages over previous efforts at this.²⁰ First, our sample size is very large, containing a large percentage of all households in Denmark. Second, we have detailed balance sheet information for not only households within our sample, but also for those excluded from our sample. Furthermore, we are able to identify interest rate risk and nominal positions held by firms, foreigners and government so that the aggregate position is zero, as required in equilibrium. This allows us to avoid some of the more egregious assumptions used in aggregating household data.

5.1 Distribution of MPX across NNP, URE and Income

The redistribution effects of monetary policy depend crucially on two household characteristics, their Net Nominal Position (henceforth NNP) and Unhedged Interest Rate Exposure (henceforth URE).

- NNP is the net value of a household’s nominal assets and liabilities. It’s relevance for analyzing the redistributive effects of monetary policy comes from the fact that an unexpected rise in the price level will decrease the wealth of households with positive nominal assets, redistributing it to those with negative NNP (who now have less real debt). In administrative data we are able to observe directly held nominal positions at the household level, including bank deposits and loans, bond holdings and mortgages. In aggregate the directly held aggregate NNP position of the household sector is negative, which from the national accounts we will see is balanced by the financial sector as well as foreigners.
- URE measures the total amount that a household plans to save at the going interest rate that period. It is the difference between all maturing assets (including income) and liabilities (including planned consumption). For example, a household with a large variable rate mortgage will likely have very negative URE. For them the entire value of their mortgage will be adjusted to the new rate. When the interest rate rises for one period they will see their disposable income (after mortgage payments) go down, and hence if they have a high MPX their spending will also decrease. To calculate URE we assume all bank deposits and bank debt have a variable rate that changes instantaneously. For mortgage debt we directly observe the amount resetting over the following year. In Denmark mortgages tend to reset only in January and July,²¹ so we assume that the new rate will only apply for half

²⁰As well as Auclert (2017), a new version of Fagereng, Holm, and Natvik (2016) also attempts to estimate these statistics.

²¹See appendix C for more details on the Danish mortgage market.

the year. For all other assets and liabilities we assume a maturity of five years. As with NNP we find households on aggregate have a negative URE position in our data and this is counterbalanced by the interest rate position of the financial sector. See appendix ***** for more detail on how we calculate NNP and URE positions.

Figure 8 show how the MPX varies across household values for URE, NNP and Income. In each case the value on the x-axis has been divided by the mean level of consumption. As mortgages in Denmark are a mixture of fixed and variable rates (see appendix C for details on the Danish mortgage market), we can think of a typical household with negative URE or NNP as having a large mortgage, while those with positive URE or NNP are wealthy households with lots of liquid wealth. A clear pattern emerges that is not evident in previous attempts to measure the distribution of MPX across these dimensions. First, and most importantly for the theory, the average MPX for those with negative URE and NNP positions is significantly greater than those with positive URE or NNP. This confirms the intuition that households who owe a lot of floating rate debt have higher MPXs than those who own this debt, and leads to an interest rate exposure channel in which lowering interest rates increases expenditure. Second, the highest MPXs don't belong to houses with large debts but instead are a characteristic of households with close to zero position in URE and NNP. These households can be characterized as being poor hand-to-mouth as they have few if any assets and their consumption behavior is not affected by changes interest rates directly.²² Note the mean levels of both URE and NNP are negative for the households in our estimation sample, so even a constant (positive) MPX would result in interest rate hikes reducing their expenditure if not balanced by indirectly held exposures.²³ While the distribution of MPX by income will not be central to our analysis, figure 8 shows a clear downward trend. If the income of lower income households decreases more than that of high income households during a monetary policy contraction, then expenditure will go down by more than the mean income weighted MPX that would be the result of a representative agent model. For comparison the distribution of MPX out of permanent income shocks across these dimensions can be found in appendix D.

5.2 Theoretical Setup and Sufficient Statistics

Auclert's method is to consider individual households' consumption response to a monetary policy shock in which i) the real rate of interest changes for one period by dR , ii) the price level makes a one time change of dP and then remains at the new level, and iii) aggregate income makes a transitory change of dY . While the dynamics here are clearly stylized, and in particular lack any lag in the economy's response, we believe such a simple experiment is highly informative as to the relative sizes of each transmission channel.

²²Neither the interest rate exposure channel nor the intertemporal substitution channel will have much impact on their consumption. Monetary policy will impact their expenditure strongly through income effects.

²³In contrast, Auclert finds a mean positive URE across households. We believe the difference is partly due to the prevalence of fixed rate mortgages in the US, but also due to underreporting of expenditures, especially in the PSID data.




Figure 8 MPX Distribution by URE, NNP and Income

Auclert divides the effect of monetary on aggregate consumption into five distinct channels:

$$\frac{dC}{C} = \underbrace{\mathcal{M} \frac{dY}{Y}}_{\text{Aggregate Income Channel}} + \underbrace{\gamma \mathcal{E}_Y \frac{dY}{Y}}_{\text{Earning's Heterogeneity Channel}} + \underbrace{-\mathcal{E}_P \frac{dP}{P}}_{\text{Fisher Channel}} + \underbrace{\mathcal{E}_R \frac{dR}{R}}_{\text{Interest Rate Exposure Channel}} + \underbrace{-\sigma \mathcal{S} \frac{dR}{R}}_{\text{Intertemporal Substitution Channel}} \quad (8)$$

where σ is the intertemporal elasticity of substitution, γ is the elasticity of relative income to aggregate income²⁴ and the five sufficient statistics, \mathcal{M} , \mathcal{E}_Y , \mathcal{E}_P , \mathcal{E}_R and \mathcal{S} are measurable in the data and defined in table 3. We chose to define these statistics to include the consumption effects coming from exposures not directly held by households. We allocate the aggregate URE and NNP exposure from our estimation sample into seven bins so that the total exposure across the economy is zero. These bins include households with (i) young (<30) and (ii) old (>55) heads, and exposures held by households indirectly through (iii) pensions funds, (iv) government, (v) non-financial corporates, (vi) financials and (vii) exposures held by the rest of the world. Within each of these bins we assume no heterogeneity so that the MPX with respect to these exposures is constant. This is a conservative assumption, likely to underestimate the size of the heterogeneous agent channels. Our assumptions on the level of these MPX's can be seen in table 4.

We define \mathcal{E}_R as:

$$\mathcal{E}_R = \frac{1}{C} \left[\sum_{i \in \text{URE deciles}} \text{MPX}_i \text{URE}_i + \sum_{j \in \text{bins}} \text{MPX}_j \text{URE}_j \right] \quad (9)$$

where i sums over the ten deciles of URE, j over the seven bins defined above and C is aggregate household expenditure in the economy. This method of dealing with the fact that aggregate exposure does not equal zero in the estimation sample is different to the approach taken by Auclert. He assumes the residual exposure is distributed equally across households in the sample. By making use of the national accounts we believe we are able to get a better handle on the likely MPX's to attach to this residual exposure. Table 3 shows the definitions we use for each of the five measurable statistics in equation 8.

²⁴Here we are making the simplifying assumptions that these quantities are common for all households, see Auclert (2017) for a discussion.

Table 3 Sufficient Statistics Definitions

Statistic	Definition	Description
\mathcal{M}	$\frac{1}{C} \left[\sum_{i \in \text{Income deciles}} \text{MPX}_i Y_i + \sum_{j \in \{\text{young, old}\}} \text{MPX}_j Y_j \right]$	Income-weighted MPX
\mathcal{E}_Y	$\mathcal{M} - \overline{\text{MPX}} \frac{Y}{C}$	Redistribution elasticity for Y
\mathcal{E}_P	$\frac{1}{C} \left[\sum_{i \in \text{NNP deciles}} \text{MPX}_i \text{NNP}_i + \sum_{j \in \text{bins}} \text{MPX}_j \text{NNP}_j \right]$	Redistribution elasticity for P
\mathcal{E}_R	$\frac{1}{C} \left[\sum_{i \in \text{URE deciles}} \text{MPX}_i \text{URE}_i + \sum_{j \in \text{bins}} \text{MPX}_j \text{URE}_j \right]$	Redistribution elasticity for R
\mathcal{S}	$1 - \frac{1}{C} \left[\sum_{i \in \text{Consumption deciles}} \text{MPX}_i C_i + \sum_{j \in \{\text{young, old}\}} \text{MPX}_j C_j \right]$	Hicksian scaling factor

Note: $\overline{\text{MPX}}$ is the mean MPX over all households in the economy. Y and C are aggregate household income and consumption respectively. bins refers to the seven categories for which we have allocated URE and NNP exposures outside our estimation sample. {young,old} are the two bins that contain young and old households (the other five bins only relevant for URE and NNP exposures as Y and C measure household income and consumption).

5.3 Out of Sample MPX

The assumptions we make about the MPX of the young and the old, as well as out of indirectly held URE and NNP exposures are shown in table 4. In each case we believe we have made conservative choices that will underestimate the size of the heterogeneous agent channels of monetary policy. For the young we choose an MPX of 0.5, in line with the rest of the population. As the young have aggregate negative exposures, choosing an MPX on the low side is conservative. Similarly for the old we choose an MPX of 0.2, on the low side, although the aggregate exposure for this group is close to zero. The assumption that there is no heterogeneity in MPX within these groups is also a very conservative assumption.

Table 4 Aggregating Redistribution Elasticities

Group	Total URE (bn Kr)	MPC	\mathcal{E}_R component
Our sample (head age 35-55)	-338	See Distribution	-0.42
Head < 30	-38	0.5	-0.02
Head > 55	-10	0.2	0.00
Pension Funds	143	0.1	0.02
Government	-120	0	0.00
Non-financial Corporate	-66	0.1	-0.01
Financial Sector	380	0.1	0.05
Rest of World	45	0	0.00
Total	0		-0.40

6 Models and Simulation

What can these results tell us about the types of models typically used for modeling consumption behavior? Two natural questions arise:

- How well does the econometric method work on standard incomplete market models?
- Can we calibrate a standard model to match the results we see, and if not are there changes we could make to the model that might help?

In this section we first introduce what is now the canonical buffer-stock savings model. We show that the econometric method is downward biased for the transitory consumption elasticity and upward biased for the permanent elasticity. When calibrated to the levels of transitory elasticity in the data, this bias is very small. Motivated by the fact that, in the data, the consumption response to transitory shocks appears larger than much of the rest of the literature, we build a model in which there is a reverse causality problem. We show that while the reverse causality can help to generate the high numbers we see for households with few liquid assets, the puzzle remains for households rich in liquid assets.

6.1 The Standard Incomplete Markets Model

Our baseline model is the now very familiar buffer-stock saving model of Carroll (1997). Given market resources (\mathbf{m}_t), households in this model maximize expected utility:

$$\mathbb{E}_t \sum_{i=t}^{\infty} \beta^i u(\mathbf{c}_i)$$

subject to the constraints:

$$\begin{aligned} \mathbf{a}_t &= \mathbf{m}_t - \mathbf{c}_t \\ \mathbf{b}_t &= R\mathbf{a}_t \\ \mathbf{y}_t &= \theta_t \mathbf{p}_t \\ \mathbf{p}_t &= \Psi_t \mathbf{p}_{t-1} \\ \mathbf{m}_t &= \mathbf{b}_t + \mathbf{y}_t \end{aligned}$$

Where the felicity function, $u(\mathbf{c})$ is CRRA. We will use this model to test the accuracy of our empirical method on simulated data. We chose two calibrations, presented in table *****. The first matches the distribution of liquid wealth in the economy, while the second is calibrated to more closely match the transitory consumption elasticity as measured in the data. In the first calibration, the economy is made up of agents with heterogenous discount factors following Carroll, Slacalek, Tokuoka, and White (2016). Agent i has a discount factor β_i where β_i is i.i.d across agents and follows a uniform distribution between β_{low} and β_{high} . These two parameters allow us to match the fact that while the mean level of liquid assets is high, about half of all households have close

to zero liquid assets. Matching the lower part of this distribution is critical to generate transitory consumption elasticities substantially above zero. In the second calibration all the agents in the economy have the same discount rate. This is chosen to give a much higher transitory consumption elasticity, close to what we see in the data. To generate such a high elasticity it is necessary to introduce an unfeasibly low discount factor. In the following sections we will show how a model of this type could possibly be reconciled with the data. Table 5 shows the estimates of ψ using simulated data. The left hand panel shows the calibration to match the distribution of liquid assets. Here we see that using $n1 = 3$ and $n2 = 5$ underestimates the true elasticity, giving a value of 0.34 versus the ‘true’ value of 0.38 obtained as $n1$ gets large (and hence the assumption that the consumption response lasts no longer than $n1 - 1$ years holds true). Thus, if this model was an accurate representation of the economy we would need to use a larger value of $n1$ or risk underestimating ψ . In the data we find our estimate for ψ is in fact much *larger* than implied by this model. In the second calibration the simulated estimate using $n1 = 3$ and $n2 = 5$, at 0.78, is very close to the ‘true’ value of 0.79.

		1	2	3	4	5	6	7	8	9	10			1	2	3	4	5	6	7	8	9	10
	n_1	1	0.21	0.23	0.23	0.24	0.24	0.25	0.25	0.25	0.25	1		1	0.62	0.64	0.66	0.66	0.67	0.67	0.67	0.67	0.67
		2		0.29	0.30	0.31	0.31	0.32	0.32	0.32	0.32	2		2		0.76	0.77	0.77	0.77	0.77	0.77	0.77	0.77
		3			0.33	0.34	0.34	0.34	0.35	0.35	0.35	3		3			0.78	0.78	0.78	0.78	0.78	0.78	0.78
		4				0.35	0.35	0.36	0.36	0.36	0.36	4		4				0.78	0.78	0.78	0.78	0.78	0.78
		5					0.36	0.36	0.36	0.36	0.37	5		5				0.79	0.78	0.78	0.79	0.79	0.79
		6						0.37	0.37	0.37	0.37	6		6					0.78	0.78	0.79	0.79	0.79
		7							0.37	0.37	0.37	7		7						0.79	0.79	0.79	0.79
		8								0.37	0.38	8		8							0.79	0.79	0.79
		9									0.38	9		9								0.78	0.78
		10										10		10									

Table 5 Simulation estimates of ψ

Table 6 shows how the estimate of ψ varies when using different values of $n1$ and $n2$ in the data. The estimates around $n1=3$ and $n2=5$ are relatively stable at 0.76, but grow significantly as $n1$ increases beyond 5. This is likely due to the permanent shock in fact having some slow reversion to the mean that makes the method unreliable for larger values of $n1$. Appendix ***** shows that when permanent income follows an AR(1) process with persistence 0.95, the method is accurate for small values of $n1$ but becomes significantly upward biased as $n1$ gets large. This creates a tension between choosing $n1$ low enough to avoid this bias, but high enough such that we allow time for the consumption response to die and do not underestimate the coefficient.

		1	2	3	4	5	6	7	8	9	10
	n_1	1	0.70	0.72	0.72	0.72	0.73	0.73	0.74	0.74	0.75
		2		0.75	0.75	0.76	0.76	0.77	0.78	0.79	0.80
		3			0.75	0.76	0.77	0.78	0.80	0.82	0.83
		4				0.77	0.78	0.80	0.82	0.85	0.87
		5					0.79	0.82	0.85	0.89	0.92
		6						0.85	0.90	0.94	0.99
		7							0.96	1.02	1.09
		8								1.12	1.26
		9									1.67
		10									

Table 6 Empirical Estimates of ψ

Table 7 shows the results of the simulation for the consumption elasticity to the permanent shock. Here we see very reliable estimates for any choice of n_1 and n_2 we care to make. Agents in the model adjust permanent consumption one for one with permanent income (in the medium run) and this comes through clearly in the estimates. The empirical counterpart to this, table 8, also shows robustness to the choice of n_1 and n_2 , apart from for larger values of n_1 . This is likely due to the same reason as the upward bias of ψ for large n_1 , namely that the permanent shock is not entirely permanent.

	1	2	3	n_2 4	5	6	7	8	9	10		1	2	3	n_2 4	5	6	7	8	9	10
n_1	1	1.12	1.07	1.04	1.03	1.02	1.01	1.01	1.00	1.00	n_1	1	1.19	1.10	1.07	1.05	1.04	1.04	1.03	1.03	1.02
	2		1.02	1.01	1.00	0.99	0.99	0.99	0.99	0.98		2		1.02	1.01	1.01	1.00	1.00	1.00	1.00	1.00
	3			0.99	0.99	0.98	0.98	0.98	0.98	0.98		3			1.00	1.00	1.00	1.00	1.00	1.00	1.00
	4				0.98	0.98	0.98	0.98	0.98	0.98		4				1.00	1.00	1.00	1.00	1.00	1.00
	5					0.98	0.98	0.98	0.98	0.98		5					1.00	1.00	1.00	1.00	1.00
	6						0.97	0.97	0.97	0.97		6						1.00	1.00	1.00	1.00
	7							0.98	0.98	0.98		7							1.00	1.00	1.00
	8								0.98	0.98		8								1.00	1.00
	9									0.97		9									1.00
	10											10									

Table 7 Simulation Estimates of ϕ

	1	2	3	n_2 4	5	6	7	8	9	10
n_1	1	0.89	0.89	0.89	0.89	0.89	0.89	0.88	0.87	0.85
	2		0.89	0.89	0.89	0.89	0.89	0.88	0.86	0.84
	3			0.90	0.90	0.90	0.89	0.87	0.85	0.84
	4				0.89	0.90	0.89	0.87	0.84	0.83
	5					0.90	0.88	0.85	0.83	0.81
	6						0.86	0.83	0.80	0.79
	7							0.80	0.78	0.77
	8								0.76	0.75
	9									0.75
	10									

Table 8 Empirical estimates of ϕ

6.2 Preference Shock Model with Elastic Labor Supply

The empirical results of this paper suggest that the consumption response to a transitory shock to income lies at the high end of the existing literature. In this section we try to reconcile the results we get with the existing literature by asking if we can build a model in which our empirical method would estimate much larger elasticities than more traditional approaches to measuring the marginal propensity to consume out of transitory shocks. The intuition we build upon is that the households may wish to work longer hours, and hence earn more, in years when their expenditure is particularly high. If this were the case the income process would not be well modeled as being exogenous and our method would have a reverse causality problem embedded in it. Here we build such a model and see how much reverse causality can plausibly contribute to the results we obtain. The model extends the standard incomplete markets model from section 6.1, incorporating both preference shocks, so that households have some years when their

utility of consumption is greater than others, and labor elasticity, so that households can adjust their income based on the marginal utility of consumption. The household's problem is to maximize expected lifetime utility:

$$\mathbb{E}_t \sum_{n=t}^{\infty} \beta^n \left(\mathcal{X}_n \frac{c_n^{1-\rho}}{1-\rho} - \frac{\ell_n^{1+\frac{1}{\xi}}}{1+\frac{1}{\xi}} \right)$$

subject to the constraints:

$$\mathbf{a}_t = \mathbf{m}_t - \mathbf{c}_t$$

$$\mathbf{b}_t = R\mathbf{a}_t$$

$$\mathbf{y}_t = l_t w_t$$

$$\ell_t = l_t \mathbf{p}_t^{\frac{1-\rho}{1+\frac{1}{\xi}}}$$

$$w_t = \theta_t \mathbf{p}_t$$

$$\mathbf{p}_t = \Psi_t \mathbf{p}_{t-1}$$

$$\mathbf{m}_t = \mathbf{b}_t + \mathbf{y}_t$$

The normalization of labor ($\ell_t = l_t \mathbf{p}_t^{\frac{1-\rho}{1+\frac{1}{\xi}}}$) is set up to allow labor supply to move elastically with transitory income, but the long run supply of labor does not depend on permanent income (as observed in the consistency of hours worked over long time periods and across countries). The key additional features of this model are i) the preference shock factor and ii) the elasticity of labor.

The preference shock factor, \mathcal{X}_t , multiplies the marginal utility of consumption in that period. When it is high the effective discount factor for consumption between this period and next is low so households will wish to spend much of their resources this period and will have little incentive to save for the next period. Equivalently when it is low, households will prefer to delay consumption to the following periods. There is little empirical evidence on the size of preference shocks, largely because consumption data usually has so much measurement error that estimating the standard deviation of consumption growth is very difficult. While a model with no preference shocks will normally imply a consumption path that is significantly smoother than the income process, our data suggests the standard deviation of income growth is 0.12 while the standard deviation of consumption growth is 0.37. Such a high standard deviation in consumption growth is likely to be mostly driven by measurement error on the asset side of the balance sheet, but it is certainly consistent with very large preference shocks. In our simulations we consider annualized preference shocks with standard deviations up to 0.5.

Labor elasticity is controlled by the Frisch elasticity ξ . When the wage (relative to permanent income) increases by $x\%$, hours worked increase by $\xi\%$. Estimates of the Frisch elasticity in micro-data studies range from 0 to 0.5, while macroeconomic studies generally find a much larger elasticity of between 2 and 4 (see Peterman (2016)). We will study a range of elasticities between 0 and 1, easily covering the microeconomic estimate

range. We will not consider estimates of the Frisch elasticity in the macroeconomic range as it seems likely to us that these estimates are high due to labor market frictions over the business cycle, rather than genuine labor supply choices of households.

	β	Frisch Elasticity						σ_q	Frisch Elasticity				
		0.00	0.13	0.25	0.38	0.50			0.00	0.13	0.25	0.38	0.50
Preference shock	0.00	0.99	0.99	0.99	0.99	0.99	Preference shock	0.00	0.07	0.07	0.06	0.05	0.05
	0.10	0.99	0.99	0.99	0.99	0.99		0.10	0.07	0.07	0.06	0.06	0.05
	0.20	0.98	0.98	0.98	0.98	0.98		0.20	0.07	0.07	0.06	0.06	0.05
	0.30	0.97	0.98	0.98	0.98	0.98		0.30	0.07	0.07	0.06	0.05	0.05
	0.40	0.96	0.96	0.97	0.97	0.98		0.40	0.07	0.07	0.06	0.05	0.04

Table 9 Fitted discount factors and transitory shock standard deviation

	ϕ	Frisch Elasticity						Std($\Delta \log c$)	Frisch Elasticity				
		0.00	0.13	0.25	0.38	0.50			0.00	0.13	0.25	0.38	0.50
Preference shock	0.00	1.00	1.00	0.99	0.99	0.99	Preference shock	0.00	0.05	0.05	0.05	0.05	0.05
	0.10	1.00	1.00	1.00	1.00	0.99		0.10	0.08	0.08	0.08	0.08	0.08
	0.20	1.02	1.01	1.01	1.00	1.00		0.20	0.13	0.13	0.13	0.13	0.13
	0.30	1.03	1.02	1.01	1.00	1.00		0.30	0.18	0.18	0.18	0.18	0.19
	0.40	1.03	1.02	1.01	0.99	0.98		0.40	0.23	0.23	0.24	0.24	0.24

Table 10 Simulation estimates of ϕ and consumption growth standard deviation

	MPC	Frisch Elasticity						ψ	Frisch Elasticity				
		0.00	0.13	0.25	0.38	0.50			0.00	0.13	0.25	0.38	0.50
Preference shock	0.00	0.17	0.13	0.11	0.09	0.08	Preference shock	0.00	0.07	0.05	0.04	0.04	0.03
	0.10	0.19	0.15	0.13	0.11	0.09		0.10	0.09	0.07	0.06	0.05	0.05
	0.20	0.25	0.20	0.16	0.13	0.11		0.20	0.14	0.12	0.11	0.12	0.13
	0.30	0.32	0.26	0.21	0.17	0.14		0.30	0.20	0.20	0.23	0.28	0.32
	0.40	0.38	0.31	0.25	0.20	0.16		0.40	0.27	0.33	0.43	0.54	0.64

Table 11 Simulation estimates of 6 month MPC and ψ

In tables 9, 10 and 11 we have varied the size of the Frisch elasticity and annualized preference shock. In each cell we have kept constant the overall annualized income growth variance and the median liquid asset to annual income ratio (equal to 0.2 in the data). To achieve this we vary the discount factor and the variance of transitory wage shocks.

Table 9 shows how the discount factor, β and the annualized transitory shock standard deviation vary. As the size of the preference shocks increase, so does the precautionary motive for households. As we have fixed the median amount of precautionary savings, the discount factor drops significantly to compensate. This effect is most prominent when there is no labor elasticity. As labor elasticity increases households are able to insure themselves against preference shocks with their labor supply, so the change in discount rate is less pronounced. The right hand panel shows the standard deviation of transitory shocks required to match the overall level of income growth variance goes down as labor supply elasticity increases. This is as expected - when the transitory wage is low households will work fewer hours. This amplifies the variance of the transitory income shock relative to the wage shock. The size of the preference shocks have little effect on the imputed size of the transitory shocks, except when both the preference shock and the labor elasticity are large. At this point the preference shocks themselves

induce significant changes in hours worked and hence income, requiring less exogenous variance in income to match the total income variance target.

The left hand panel of table 10 shows the estimate of ϕ (the consumption elasticity to permanent shocks) is close to 1 for variations of preference shocks and labor elasticities. This is unsurprising as labor does not respond to a change in permanent income. The right hand panel shows a very significant increase in the standard deviation of consumption growth as the size of the preference shocks increases. With no preference shocks, the standard deviation of consumption growth (0.05) is about half of the standard deviation of income (0.12). As the size of preference shocks increases, so does consumption growth variance, with the standard deviation growing to 0.23 for large preference shocks. This is still much smaller than 0.37, which comes directly from the data, although this high number from the data is likely to be contaminated with measurement error in assets. A further consideration is that much of the observed variance in expenditure growth will be due to durable items, such as home improvements and vehicles. We analyze the effect of durables on our estimates in section 6.6, but to the extent that these goods can be financed, our model with no borrowing may overestimate both the expenditure variance and the labor supply response to preference shocks.

Table 11 compares the actual mean six month MPC in the model with our empirical method for estimating the transitory expenditure elasticity. While these concepts are clearly not equivalent, comparing the two gives a good understanding as to what might be going on. The left hand panel shows that both preference shocks and labor elasticity, often both missing in consumption models for simplicity, have quantitatively significant impacts on the implied marginal propensity to consume. Increasing the Frisch elasticity from 0 to 0.5 (the full range of micro-estimates) decreases the six month MPC from 17% to just 8%. This is because households now have an extra tool with which to insure against low consumption. When they receive a negative transitory shock to their wealth, they will consume less, which in turn will increase their marginal utility of consumption and induce them to work more hours. Therefore their actual income loss will be lower than the shock to their wealth and they will reduce their consumption by less than if they were unable to adjust their labor supply. In contrast, increasing the size of the preference shocks greatly increases the marginal propensity to consume. This is a result of the higher precautionary savings motive and consequently lower discount factor, even while median savings are unchanged. Households with a large positive preference shock will have very little motive to save, as with such low level of savings they will soon get close to the natural borrowing constraint with an MPC close to unity. Many recent papers, such as Krueger, Mitman, and Perri (2016), have attempted to carefully quantify the macroeconomic dynamic consequences of a serious heterogeneous agent model, but thus far have not included significant preference shocks in their calibrations. The evidence here suggests that such shocks may have a quantitatively important role to play, especially in increasing the marginal propensity to consume. To the extent that the precautionary motive is driven by preference shocks as opposed to income shocks, social insurance for unemployment will not reduce precautionary savings as much as these models presently suggest.

The right hand panel of table 11 shows the effect of preference shocks and labor

elasticity on our empirical estimates of ψ , the transitory expenditure elasticity. The top row shows that our estimate is lower than the six month MPC (due to the fact that at these low levels of MPC, more than three years is required for the transitory effect to decay away). It does however follow the same pattern as the MPC and falls in magnitude as the ability of households to adjust labor supply increases. Similarly, going down the first row shows that the estimated expenditure elasticity increases with the preference shock. However, the similarity to the MPC table ends when we increase *both* labor elasticity *and* the size of the preference shocks. Our estimate can grow large, up to a value of 0.64, getting close to our empirical estimates, when the Frisch elasticity is 0.5 and the preference shock standard deviation is 0.4. This measured consumption elasticity to transitory income shocks now bears little relation to the six month MPC (which is 0.16). Instead it is being driven by reverse causality, whereby preference shocks are driving consumption along with the decision to increase labor. The observed ‘shocks’ to income are therefore highly correlated with consumption, but they are not causing the consumption dynamics exogenously.

6.3 Robustness to Misspecification

In general the estimated permanent and transitory shock variance will be sensitive to the exact assumptions made about the income process. For example allowing the permanent component to mean revert (as in Ahn, Kaplan, Moll, Winberry, and Wolf (2017) or Arellano, Blundell, and Bonhomme (2017)) or allowing for correlation between the permanent and transitory shocks will affect how the total observed variance is divided between permanent and transitory. However, as the estimates we obtain for transitory and permanent MPX are close to equal to each other these results are relatively robust to different specifications of the income process.²⁵

A key part of our identification method is the assumption that the impulse response for both income and consumption to a transitory shock has decayed to zero two years after the initial shock. We used growth over 3, 4 and 5 years for our identification because three years is the minimum to allow for generic impulse responses up to two years and demanding that households have more than five years of continuous data reduces the sample and possibly introduces sample selection problems. Figure 9 makes use of the fact that our model is overidentified with the rich data available to us, so we can test some of this assumption. The solid black line shows the empirically measured variance of income growth over 1 to 7 years. The solid red line draws the straight line equation 4 with transitory and permanent variance parameters estimated using minimum distance on 3, 4 and 5 years of income growth. The empirical variance of income growth over two and three years is lower than this red line showing equation 4 is not valid for $n < 3$, consistent with some persistence in the transitory shock extending out to two years. However, the empirical variance over 6 and 7 years of income growth falls almost exactly on the red line implied by the model. This is strong evidence that the transitory

²⁵We plan to carry out a number of robustness tests by changing the assumptions made about both the income and consumption processes. Our initial results suggest the estimated variances can change significantly, but the estimated MPX numbers are relatively robust.




Figure 9 Variance and Covariance Over n Years of Growth

component and income does not persist more than two years and that the permanent component does not mean revert at any significant rate.²⁶

The dashed black line and solid green line repeat this experiment with the covariance of income and consumption over increasing years of growth. This time the solid green line is equation 6 with parameters identified by 3, 4 and 5 years of growth. As with income, the model says that this covariance should grow linearly for $n \geq 3$. The assumption that the transitory consumption response goes to zero within 2 years is less common than the equivalent assumption on income and is key to our identification. The data again show that for $n < 3$ the covariance is below the green line, but that for growth over 6 and 7 years the model does a good job.

6.4 Why is the MPX out of Transitory Shocks so High?

The results of this paper for permanent shocks is generally in line with both theory and previous empirical work. However, the MPX out of transitory shocks is at the high end of empirical work and is especially high relative to standard consumption theory.

Table 1 shows that our estimate of the MPX out of transitory shocks is at the high end, but not extreme. Our measure includes all spending on both durables and non-

²⁶We plan to carry out formal overidentification tests here. As we have so much data it is very likely that such a test suggest our model is misspecified. This would not be a concern (all models are simplifications) unless it is economically significant, something figure 9 shows is not the case.

durables and represents the elasticity of consumption with transitory income received over a full year, a longer period than for many of the estimates in table 1. Estimates from Agarwal and Qian (2014), Parker, Souleles, Johnson, and McClelland (2013) and Souleles (1999) are all similar or higher than ours. However, Fagereng, Holm, and Natvik (2016) estimate much lower MPXs (around 0.4) using well identified lottery winnings in a country (Norway) that is likely more similar to Denmark than the US. Furthermore, the expenditure data they use is imputed in a similar way to ours, so in order to take our results seriously we should have some idea why they may be so different to those from this paper. One answer may be that household’s consumption response to a lottery win is different to other types of transitory shocks they may receive. While our initial impression is that households would spend more out of a lottery win (especially if they spend on celebrating the win), it is also possible that households put their winnings into a separate ‘mental account’ (Thaler (1985)) from which they spend less than from their labor earnings.

Another possibility, that deserves further investigation, is that our method may be picking up reverse causality between income and consumption. In our model, as is common in this literature, we treat the income process as exogenous. However, households may be able to adjust their labor supply in response to shocks. In particular, in years in which households have large spending needs (for example if they buy a car, carry out home improvements or pay for an offspring’s wedding) it may be the case that they simultaneously increase their labor supply and hence earnings. Even (or especially) if the labor supply response is small, our model would see that as a very large increase in spending in response to a small increase in income. In future work we plan to address this in a quantitative model in which there are both transitory shocks to income and transitory shocks to consumption preferences. It is unclear to us at the moment how quantitatively important such a mechanism may be. Some empirical evidence that this mechanism may be important is that in households with two earners, the MPX out of transitory shocks from the secondary earner is significantly higher than for the primary earner. It is well established that the labor elasticity of secondary earners is higher than for primary earners so the possibility of reverse causality may be stronger for them.

A final possibility is that measurement error is driving our results to some extent. Our method is robust to classical measurement error in consumption. The permanent MPX is robust to classical measurement error in income while the transitory MPX would be biased downwards with classical measurement error in income uncorrelated with consumption. However, our method of imputing expenditure from income and wealth will pass any measurement error in income over to expenditure. Thus the portion of transitory income variance that is due to measurement error will have a related MPX of 1. Due to the fact that our income data comes from third party reported tax data we believe this portion is likely very small. We plan to investigate this by reducing the sample to households for whom we have reason to believe our income data is particularly accurate. A further test we plan to carry out is to use car purchases instead of total expenditure. Register data on car purchases are available from the Danish car registry. This gives us a source of very well measured expenditure totally independent of our imputed expenditure method. This will act as a robustness test to our surprising result

that even households with high levels of liquid wealth have expenditure tied closely with income.

6.5 Robustness to Auto Regressive Permanent Income

The previous sections attempted to reconcile the results we obtain with modern consumption savings models. Our conclusion from that exercise is that these models cannot reproduce the observed results unless we assume empirically unrealistic levels of preference shocks and income elasticity. In the models we considered we did not change the income process, which we assumed consisted of both a permanent and a transitory component. In this section we will consider non-structural models, closer to that of our empirical method itself, but instead explore how misspecification in the income process may affect our results. In particular we will simulate a model of the form:

$$\begin{aligned} p_t &= \rho p_{t-1} + \varepsilon_t \\ y_t &= p_t + q_t \\ c_t &= \phi y_t + \psi q_t \end{aligned}$$

where p_t is the log of permanent income, y_t the log of total income, and q_t a transitory shock to income. We choose the variance of the permanent and transitory shocks to be equal. We also set the consumption elasticity to permanent shocks, ϕ to be 0.8 (to match our empirical results) and the consumption elasticity to transitory shocks, ψ , to be 0.4. This is chosen to be closer to the rest of the empirical literature on marginal propensities to consume. We are interested in seeing if our empirical method may significantly bias this 0.4 value in the case that the income process is misspecified.

ϕ	1	2	n_1 3	4	5	6	7	8	ψ	1	2	n_1 3	4	5	6	7	8
ρ 1.0	0.80	0.80	0.80	0.80	0.80	0.80	0.80	0.80	1.0	0.40	0.40	0.40	0.40	0.40	0.40	0.40	0.40
0.98	0.80	0.80	0.80	0.80	0.80	0.80	0.80	0.80	ρ 0.98	0.41	0.42	0.43	0.44	0.46	0.47	0.49	0.51
0.96	0.80	0.80	0.80	0.80	0.80	0.80	0.80	0.80	0.96	0.41	0.43	0.45	0.47	0.50	0.52	0.54	0.56
0.94	0.80	0.80	0.80	0.80	0.80	0.80	0.80	0.80	0.94	0.42	0.44	0.47	0.50	0.52	0.55	0.57	0.59
0.92	0.80	0.80	0.80	0.80	0.80	0.80	0.80	0.80	0.92	0.42	0.45	0.49	0.52	0.55	0.57	0.59	0.61
0.9	0.80	0.80	0.80	0.80	0.80	0.80	0.80	0.80	0.9	0.43	0.46	0.50	0.53	0.56	0.59	0.61	0.62

Table 12 Estimates with AR(1) income process

Table 12 shows the results of our simulation for various values of ρ and n_1 (n_2 is always chosen to be $n_1 + 2$). Remember in our core results we use values of $n_1 = 3$ and $n_2 = 5$. The left hand panel of table 12 shows the misspecification of the income process does not bias the estimate of ϕ . However, the right hand panel shows a slight upward bias for our estimate of ψ as the value of ρ decreases. This bias increases with the value of n_1 . This is because the total income variance does not grow linearly with time. Instead the growth in variance starts to taper off as the period of over which income growth is measured increases. This leads to the empirical methodology interpreting much of the ‘permanent’ variance as transitory, and associating a higher consumption elasticity with

it. This effect is small for reasonable values of ρ at $n_1 = 3$. Furthermore as the difference between the true ϕ and ψ becomes smaller, so does this bias. When $\phi = \psi$ there is no bias at all, as it makes no difference if the empirical methodology interprets the income variance as permanent or transitory.

6.6 Durables

A further critique of our empirical methodology is that it does not take account of durable goods, while our data includes all spending (except on real estate) and therefore includes large and durable goods such as cars and home improvements. The empirical model assumes that in response to a transitive income shock, expenditure increases temporarily for up to two years. This is entirely consistent with a model that includes durable goods. However, the model assumes that in response to a permanent shock to income, expenditure increases once to a new permanent level. A model that included permanent goods would instead imply a large one off expenditure on durable goods to get the household up to their desired stream of durable good services, followed by a decrease back to a permanent level of spending that accounts for replenishing the higher level of depreciating durable goods.

To make this idea more explicit, it will help to write down a simple model. The model will show that our empirical methodology continues to estimate the consumption elasticities to permanent and transitory shocks, but that these need to be interpreted carefully. The model uses the same income process as section ?? and takes the same shortcuts, working in logs rather than levels. The same ideas as in appendix A can be used to formalize this model in levels. Remembering the income process is made up of two martingale processes, P_t and Q_t , which may have jumps. Instantaneous income is given by:

$$dy_t = \left(\int_0^t dP_s \right) dt + dQ_t$$

while instantaneous expenditure now has both a durable and non-durable component:

$$dc_t = \phi_{nd} \left(\int_0^t dP_s \right) dt + \phi_d dP_t + \psi dQ_s$$

Here we have assumed that the expenditure response to transitory shocks is instantaneous, but it would not change things to assume as before that the response decays to zero after two years. However, it is important that the durable component of the expenditure response to permanent shocks occurs instantaneously with the shock (or very soon after). Aggregating income and consumption annually gives:

$$\begin{aligned} \Delta^N \bar{y}_T = & \left(\int_{T-N-1}^{T-N} (s - (T - N - 1)) dP_s + \int_{T-N}^{T-1} dP_s + \int_{T-1}^T (T - s) dP_s \right) \\ & + \left(\int_{T-1}^T dQ_t - \int_{T-N-1}^{T-N} dQ_t \right) \end{aligned}$$

$$\begin{aligned}
\Delta^N \bar{c}_T &= \phi_{nd} \left(\int_{T-N-1}^{T-N} (s - (T - N - 1)) dP_s + \int_{T-N}^{T-1} dP_s + \int_{T-1}^T (T - s) dP_s \right) \\
&+ \phi_d \left(\int_{T-1}^T dP_t - \int_{T-N-1}^{T-N} dP_t \right) \\
&+ \psi \left(\int_{T-1}^T dQ_t - \int_{T-N-1}^{T-N} dQ_t \right)
\end{aligned}$$

From this we can calculate the covariance:

$$\begin{aligned}
\text{Cov}(\Delta^N \bar{c}_T, \Delta^N \bar{y}_T) &= \phi_{nd} \text{Var}(\Delta^N \bar{y}_T) \\
&+ \phi_d \left(\int_{T-1}^T (T - s) \sigma_P^2 dt - \int_{T-N-1}^{T-N} (s - (T - N - 1)) \sigma_P^2 dt \right) \\
&+ \psi \left(\int_{T-1}^T \sigma_Q^2 dt + \int_{T-N-1}^{T-N} \sigma_Q^2 dt \right) \\
&= \phi_{nd} \left(n - \frac{1}{3} \right) \sigma_P^2 + 0 + 2\psi \sigma_Q^2
\end{aligned}$$

So the durable component of the covariance cancels out and our identification method correctly identifies ϕ_{nd} and ψ , but is unable to identify ϕ_d .

7 Conclusion

Our novel method of measuring the consumption response to income shocks, along with its application to Danish registry data, shows that households with little liquid wealth have spending that moves almost one for one with income, while even households with relatively high levels of liquid wealth also tie their spending closely to their income. We have shown serious flaws in previous methods that have tried to uncover this relation by imposing structure on the income and consumption processes. Our method and results have also opened up new questions, including whether it is reasonable to assume income follows exogenous shocks when doing this kind of work. In future work we plan to look more closely at how labor responds to preference shocks and what implications that might have for the macroeconomy. Our results suggest that households may use adjustments in their labor supply as much as their wealth to self-insure. In a recession voluntarily adjusting one's labor supply become much harder to do as work is much less freely available. One response to this may be for households to increase their precautionary savings to protect against future income or spending shocks. Such a mechanism would further dampen aggregate demand and aggregate the recession.

References

- ABILDGREN, KIM, ANDREAS KUCHLER, AMERICA SOLANGE, AND HENRIK SORESENSEN (2018): “Consistency between household-level consumption data from registers and surveys,” .
- AGARWAL, SUMIT, AND WENLAN QIAN (2014): “Consumption and Debt Response to Unanticipated Income Shocks: Evidence from a Natural Experiment in Singapore,” *American Economic Review*, 104(12), 4205–4230.
- AHN, SEHYOUN, GREG KAPLAN, BENJAMIN MOLL, THOMAS WINBERRY, AND CHRISTIAN WOLF (2017): “When Inequality Matters for Macro and Macro Matters for Inequality,” *NBER*.
- ARELLANO, MANUEL, RICHARD BLUNDELL, AND STÉPHANE BONHOMME (2017): “Earnings and Consumption Dynamics: A Nonlinear Panel Data Framework,” *Econometrica*, 85(3), 693–734.
- AUCLERT, ADRIEN (2017): “Monetary Policy and the Redistribution Channel,” Working Paper 23451, National Bureau of Economic Research.
- BAKER, SCOTT R. (2015): “Debt and the Consumption Response to Household Income Shocks,” SSRN Scholarly Paper ID 2541142, Social Science Research Network, Rochester, NY.
- BLUNDELL, RICHARD, LUIGI PISTAFERRI, AND IAN PRESTON (2008): “Consumption Inequality and Partial Insurance,” *American Economic Review*, 98(5), 1887–1921.
- BROWNING, MARTIN, AND M. DOLORES COLLADO (2001): “The Response of Expenditures to Anticipated Income Changes: Panel Data Estimates,” *American Economic Review*, 91(3), 681–692.
- BROWNING, MARTIN, AND SØREN LETH-PETERSEN (2003): “Imputing consumption from income and wealth information,” *Economic Journal*, 113(488), F282–F301.
- CAMPBELL, JOHN Y., AND N. GREGORY MANKIW (1989): “Consumption, Income and Interest Rates: Reinterpreting the Time Series Evidence,” NBER Chapters, National Bureau of Economic Research, Inc.
- CARROLL, CHRISTOPHER, AND ANDREW SAMWICK (1997): “The nature of precautionary wealth,” *Journal of Monetary Economics*, 40(1), 41–71.
- CARROLL, CHRISTOPHER D. (1997): “Buffer Stock Saving and the Life Cycle/Permanent Income Hypothesis,” *Quarterly Journal of Economics*, CXII(1), 1–56.
- CARROLL, CHRISTOPHER D., JIRI SLACALEK, KIICHI TOKUOKA, AND MATTHEW N. WHITE (2016): “The Distribution of Wealth and the Marginal Propensity to Consume,” Draft, Johns Hopkins University, At <http://econ.jhu.edu/people/ccarroll/papers/cstwMPC> (this is a revision of <http://econ.jhu.edu/people/ccarroll/papers/cstMPC\texttt{cstMPC}>).

- COMMAULT, JEANNE (2017): “How Does Consumption Respond to a Transitory Income Shock? Reconciling Natural Experiments and Structural Estimations,” .
- CORONADO, JULIA LYNN, JOSEPH P. LUPTON, AND LOUISE M. SHEINER (2005): “The Household Spending Response to the 2003 Tax Cut: Evidence from Survey Data,” FEDS discussion paper 32, Federal Reserve Board.
- CRAWLEY, EDMUND (2018): “Time Aggregation in Panel Data on Income and Consumption,” Mimeo, Department of Economics, Johns Hopkins University.
- DAVID, MARTIN, KENT H. MARQUIS, JEFFREY C MOORE, LINDA L. STINSON, AND EDWARD J. WELNIAK (1997): “Income Measurement Error in Surveys: A Review - Semantic Scholar,” .
- FAGERENG, ANDREAS, AND ELIN HALVORSEN (2015): “Imputing consumption from Norwegian income and wealth registry data,” Discussion Paper, Statistics Norway, Research Department.
- FAGERENG, ANDREAS, M. B. HOLM, AND G. J. NATVIK (2016): “MPC Heterogeneity and Household Balance Sheets,” *Discussion Paper, Statistics Norway*.
- FUSTER, ANDREAS, GREG KAPLAN, AND BASIT ZAFAR (2018): “What Would You Do With \$500? Spending Responses to Gains, Losses, News and Loans,” Working Paper 24386, National Bureau of Economic Research.
- GANONG, PETER, AND PASCAL NOEL (2017): “Consumer Spending During Unemployment: Positive and Normative Implications,” .
- GELMAN, MICHAEL (2016): “What Drives Heterogeneity in the Marginal Propensity to Consume? Temporary Shocks vs Persistent Characteristics,” .
- GELMAN, MICHAEL, YURIY GORODNICHENKO, SHACHAR KARIV, DMITRI KOUSTAS, MATTHEW D. SHAPIRO, DAN SILVERMAN, AND STEVEN TADELIS (2016): “The Response of Consumer Spending to Changes in Gasoline Prices,” Working Paper 22969, National Bureau of Economic Research.
- GELMAN, MICHAEL, SHACHAR KARIV, MATTHEW D. SHAPIRO, DAN SILVERMAN, AND STEVEN TADELIS (2014): “Harnessing naturally occurring data to measure the response of spending to income,” *Science*, 345(6193), 212–215.
- HAUSMAN, JOSHUA K. (2012): “Fiscal Policy and Economic Recovery: The Case of the 1936 Veterans’ Bonus,” mimeo, University of California, Berkeley.
- HEATHCOTE, JONATHAN, FABRIZIO PERRI, AND GIOVANNI L. VIOLANTE (2010): “Unequal we stand: An empirical analysis of economic inequality in the United States, 1967-2006,” *Review of Economic Dynamics*, 13(1), 15–51.
- HRYSKO, DMYTRO, AND IOURII MANOVSKII (2018): “Income Dynamics and Consumption Insurance,” .

- HSIEH, CHANG-TAI (2003): “Do Consumers React to Anticipated Income Changes? Evidence from the Alaska Permanent Fund,” *American Economic Review*, 99, 397–405.
- JAPPELLI, TULLIO, AND LUIGI PISTAFERRI (2010): “The Consumption Response to Income Changes,” *Annual Review of Economics*, 2(1), 479–506.
- (2014): “Fiscal Policy and MPC Heterogeneity,” *American Economic Journal: Macroeconomics*, 6(4), 107–136.
- JOHNSON, DAVID S., JONATHAN A. PARKER, AND NICHOLAS S. SOULELES (2006): “Household Expenditure and the Income Tax Rebates of 2001,” *American Economic Review*, 96(5), 1589–1610.
- (2009): “The Response of Consumer Spending to Rebates During an Expansion: Evidence from the 2003 Child Tax Credit,” working paper, The Wharton School.
- KAPLAN, GREG, BENJAMIN MOLL, AND GIOVANNI L. VIOLANTE (2016): “Monetary Policy According to HANK,” Working Paper 21897, National Bureau of Economic Research.
- KAPLAN, GREG, AND GIOVANNI L. VIOLANTE (2010): “How Much Consumption Insurance beyond Self-Insurance?,” *American Economic Journal: Macroeconomics*, 2(4), 53–87.
- KAPLAN, GREG, AND GIOVANNI L. VIOLANTE (2018): “Microeconomic Heterogeneity and Macroeconomic Shocks,” Working Paper 24734, National Bureau of Economic Research.
- KOIJEN, RALPH, STIJN VAN NIEUWERBURGH, AND ROINE VESTMAN (2014): “Judging the Quality of Survey Data by Comparison with "Truth" as Measured by Administrative Records: Evidence From Sweden,” in *NBER Chapters: Improving the Measurement of Consumer Expenditures*, pp. 308–346. National Bureau of Economic Research, Inc.
- KRUEGER, DIRK, KURT MITMAN, AND FABRIZIO PERRI (2016): “Macroeconomics and Household Heterogeneity,” Working Paper 22319, National Bureau of Economic Research.
- KRUSELL, PER, AND JR. SMITH, ANTHONY (1998): “Income and Wealth Heterogeneity in the Macroeconomy,” *Journal of Political Economy*, 106(5), 867–896.
- LUSARDI, ANNAMARIA (1996): “Permanent Income, Current Income, and Consumption: Evidence from Two Panel Data Sets,” *Journal of Business and Economic Statistics*, 14(1), 81–90.
- MANOVSKII, IOURII, AND DMYTRO HRYSHKO (2017): “How Much Consumption Insurance in the U.S.?,” Discussion Paper 1584, Society for Economic Dynamics.
- MEGHIR, COSTAS, AND LUIGI PISTAFERRI (2004): “Income Variance Dynamics and Heterogeneity,” *Econometrica*, 72(1), 1–32.
- MOFFITT, ROBERT A., AND PETER GOTTSCHALK (2012): “Trends in the Transitory Variance of Male Earnings Methods and Evidence,” *Journal of Human Resources*, 47(1), 204–236.
- MOFFITT, ROBERT A., AND SISI ZHANG (2018): “Income Volatility and the PSID: Past Research and New Results,” Working Paper 24390, National Bureau of Economic Research.

- NIELSEN, HELENA SKYT, AND ANNETTE VISSING-JORGENSEN (2004): *The Impact of Labor Income Risk on Educational Choices: Estimates and Implied Risk Aversion*.
- PARKER, JONATHAN A. (1999): “The Reaction of Household Consumption to Predictable Changes in Social Security Taxes,” *American Economic Review*, 89(4), 959–973.
- PARKER, JONATHAN A, NICHOLAS S SOULELES, DAVID S JOHNSON, AND ROBERT MCCLELLAND (2013): “Consumer spending and the economic stimulus payments of 2008,” *The American Economic Review*, 103(6), 2530–2553.
- PETERMAN, WILLIAM B. (2016): “Reconciling Micro and Macro Estimates of the Frisch Labor Supply Elasticity,” *Economic Inquiry*, 54(1), 100–120.
- SAHM, CLAUDIA R., MATTHEW D. SHAPIRO, AND JOEL B. SLEMROD (2010): “Household Response to the 2008 Tax Rebate: Survey Evidence and Aggregate Implications,” *Tax Policy and the Economy*, 24, 69–110.
- SHAPIRO, MATTHEW D., AND JOEL SLEMROD (2009): “Did the 2008 Tax Rebates Stimulate Spending?,” *American Economic Review*, 99(2), 374–79.
- SOULELES, NICHOLAS S. (1999): “The Response of Household Consumption to Income Tax Refunds,” *American Economic Review*, 89(4), 947–958.
- (2002): “Consumer Response to the Reagan Tax Cuts,” *Journal of Public Economics*, 85, 99–120.
- THALER, RICHARD (1985): “Mental Accounting and Consumer Choice,” *Marketing Science*, 4(3), 199–214.
- VARDARDOTTIR, ARNA, AND MICHAELA PAGEL (2016): “The Liquid Hand-to-Mouth: Evidence from a Personal Finance Management Software,” 2016 Meeting Paper 789, Society for Economic Dynamics.
- VIOLANTE, GIANLUCA, GREG KAPLAN, AND JUSTIN WEIDNER (2014): “The Wealthy Hand-to-Mouth,” 2014 Meeting Paper 192, Society for Economic Dynamics.
- WORKING, HOLBROOK (1960): “Note on the Correlation of First Differences of Averages in a Random Chain,” *Econometrica*, 28(4), 916–918.

Appendix

A Identification with Time Aggregation

In this section we formalize the continuous time model and calculate the relevant variance and covariances. We begin by defining permanent income. Let p_t for $t \in \mathbb{R}^+$ be a martingale process (possibly with jumps) with independent stationary increments and ν_p be such that $\mathbb{E}(e^{p_t - p_{t-1}}) = e^{\nu_p}$. Define the permanent component of income as:

$$P_t = e^{p_t - t\nu_p}$$

Note that $\mathbb{E}\left(\frac{P_{t+s}}{P_t}\right) = 1$ for all $s \geq 0$.

Next we define transitory income. Let q_t on $t \in \mathbb{R}^+$ also be a martingale process, independent of p_t , with independent stationary increments. Let $f : \mathbb{R}^+ \rightarrow \mathbb{R}$ be the impulse response of income to changes in q_t . We will assume that the impulse response to a transitory shock to income is over after two years, that is $f(s) = 0$ for $s > 2$. The transitory component of income is then defined as:

$$\theta_t = e^{\int_{t-2}^t f(t-s)dq_s - \nu_q}$$

where $e^{\nu_q} = \mathbb{E}e^{\int_{t-2}^t f(t-s)dq_s}$ so that $\mathbb{E}\theta_t = 1$.

We are now in a position to talk about total income. Total income *flow* at time t is given by:

$$\begin{aligned} Y_t &= P_t \theta_t \\ &= e^{p_t - t\nu_p + \int_{t-2}^t f(t-s)dq_s - \nu_q} \end{aligned}$$

Observable income is the sum of income *flow* over a one year period, that is:

$$\bar{Y}_T = \int_{T-1}^T P_t \theta_t dt$$

We will be focused on the log of observable income growth over N years:

$$\begin{aligned} \Delta^N \log(\bar{y}_T) &= \log\left(\int_{T-1}^T P_t \theta_t dt\right) - \log\left(\int_{T-N-1}^{T-N} P_t \theta_t dt\right) \\ &= \log\left(\frac{P_{T-1}}{P_{T-N}}\right) + \log\left(\int_{T-1}^T \frac{P_t}{P_{T-1}} \theta_t dt\right) - \log\left(\int_{T-N-1}^{T-N} \frac{P_t}{P_{T-N}} \theta_t dt\right) \quad (10) \end{aligned}$$

Note that if $N \geq 3$ each of the three components of equation 10 are mutually independent because both p_t and q_t have independent increments, and θ_t is independent of q_s for $s < t - 2$ and $s > t$. Defining $\mathcal{P}_{T,N}$, $\mathcal{Q}_{T,N}^1$ and $\mathcal{Q}_{T,N}^2$ to be the three parts of the sum in equation 10 respectively, we have:

$$\begin{aligned} \mathcal{P}_{T,N} &= \log\left(\frac{P_{T-1}}{P_{T-N}}\right) \\ \Rightarrow \text{Var}(\mathcal{P}_{T,N}) &= (N-1)\text{Var}\left(\log\left(\frac{P_T}{P_{T-1}}\right)\right) \end{aligned}$$

$$= (N - 1)\sigma_P^2$$

where σ_P^2 is defined to be $\text{Var}\left(\log\left(\frac{P_T}{P_{T-1}}\right)\right)$, which does not depend on T because p_t has independent increments. Moving on to the components that contain a mix of both permanent and transitory income, and defining $\bar{\theta}_T = \int_{T-1}^T \theta_t dt$, we have

$$\begin{aligned}\mathcal{Q}_{T,N}^1 &= \log\left(\int_{T-1}^T \frac{P_t}{P_{T-1}} \theta_t dt\right) \\ &= \log\left(\int_{T-1}^T \theta_t dt + \int_{T-1}^T \left(\frac{P_t}{P_{T-1}} - 1\right) \theta_t dt\right) \\ &= \log(\bar{\theta}_T) + \log\left(1 + \int_{T-1}^T \left(\frac{P_t}{P_{T-1}} - 1\right) \frac{\theta_t}{\bar{\theta}_T} dt\right) \\ &\approx \log(\bar{\theta}_T) + \int_{T-1}^T \left(\frac{P_t}{P_{T-1}} - 1\right) \frac{\theta_t}{\bar{\theta}_T} dt\end{aligned}$$

Where the approximation holds so long as $\frac{P_t}{P_{T-1}}$ is close to 1 for $T - 1 \leq t \leq T$, that is the permanent shock does not move a lot in the course of one year. Define:

$$\sigma_\theta^2 = \text{Var}\left(\log(\bar{\theta}_T)\right)$$

So that

$$\begin{aligned}\text{Var}(\mathcal{Q}_{T,N}^1) &\approx \sigma_\theta^2 + \mathbb{E}\left(\int_{T-1}^T \left(\frac{P_t}{P_{T-1}} - 1\right) \frac{\theta_t}{\bar{\theta}_T} dt\right)^2 \\ &= \sigma_\theta^2 + \mathbb{E}\left(\int_{T-1}^T \int_{T-1}^T \left(\frac{P_t}{P_{T-1}} - 1\right) \left(\frac{P_s}{P_{T-1}} - 1\right) \frac{\theta_t \theta_s}{\bar{\theta}_T^2} dt ds\right) \\ &= \sigma_\theta^2 + \int_{T-1}^T \int_{T-1}^T \mathbb{E}\left(\left(\frac{P_{\min(t,s)}}{P_{T-1}}\right)^2 \frac{P_{\max(t,s)}}{P_{\min(t,s)}} - \frac{P_t}{P_{T-1}} - \frac{P_s}{P_{T-1}} - 1\right) \mathbb{E}\left(\frac{\theta_t \theta_s}{\bar{\theta}_T^2}\right) dt ds \\ &= \sigma_\theta^2 + \int_{T-1}^T \int_{T-1}^T \text{Var}\left(\frac{P_{\min(t,s)}}{P_{T-1}}\right) \mathbb{E}\left(\frac{\theta_t \theta_s}{\bar{\theta}_T^2}\right) dt ds \\ &\approx \sigma_\theta^2 + \sigma_P^2 \int_{T-1}^T \int_{T-1}^T \min(t, s) \mathbb{E}\left(\frac{\theta_t \theta_s}{\bar{\theta}_T^2}\right) dt ds \\ &= \sigma_\theta^2 + \sigma_P^2 \int_{T-1}^T \int_{T-1}^T \min(t, s) \mathbb{E}\left(\left(1 + \frac{\theta_t - \bar{\theta}_T}{\bar{\theta}_T}\right) \left(1 + \frac{\theta_s - \bar{\theta}_T}{\bar{\theta}_T}\right)\right) dt ds \\ &= \sigma_\theta^2 + \sigma_P^2 \int_{T-1}^T \int_{T-1}^T \min(t, s) \left(1 + \mathbb{E}(\hat{\theta}_{t,T}) + \mathbb{E}(\hat{\theta}_{s,T}) + \mathbb{E}(\hat{\theta}_{t,T} \hat{\theta}_{s,T})\right) dt ds\end{aligned}$$

where $\hat{\theta}_{t,T} = \frac{\theta_t - \bar{\theta}_T}{\bar{\theta}_T}$. Continuing:

$$\begin{aligned}
\text{Var}(\mathcal{Q}_{T,N}^1) &\approx \sigma_\theta^2 + \sigma_P^2 \int_{T-1}^T \int_{T-1}^T \min(t, s) dt ds \\
&\quad + \underbrace{\sigma_P^2 \int_{T-1}^T \int_{T-1}^T \min(t, s) \left(\mathbb{E}(\hat{\theta}_{t,T}) + \mathbb{E}(\hat{\theta}_{s,T}) + \mathbb{E}(\hat{\theta}_{t,T} \hat{\theta}_{s,T}) \right) dt ds}_{\approx 0} \\
&= \sigma_\theta^2 + \sigma_P^2 \int_{T-1}^T \left(\int_{T-1}^s t dt + \int_s^T s dt \right) ds \\
&= \sigma_\theta^2 + \frac{1}{3} \sigma_P^2
\end{aligned}$$

A very similar calculation shows that:

$$\text{Var}(\mathcal{Q}_{T,N}^2) \approx \sigma_\theta^2 + \frac{1}{3} \sigma_P^2$$

So we get that:

$$\begin{aligned}
\text{Var}(\Delta^N \log(\bar{y}_T)) &= \text{Var}(\mathcal{P}_{T,N}) + \text{Var}(\mathcal{Q}_{T,N}^1) + \text{Var}(\mathcal{Q}_{T,N}^2) \\
&\approx (N-1)\sigma_P^2 + (\sigma_\theta^2 + \frac{1}{3}\sigma_P^2) + (\sigma_\theta^2 + \frac{1}{3}\sigma_P^2) \\
&= (N - \frac{1}{3})\sigma_P^2 + 2\sigma_\theta^2
\end{aligned}$$

Now we turn to consumption. Consumption responds to permanent income with elasticity ϕ , while the impulse response to a transitory shock is given by some function $g : \mathbb{R}^+ \rightarrow \mathbb{R}$ with $g(s) = 0$ for $s > 2$. Total consumption *flow* is then given by:

$$C_t = C_t^P C_t^\theta$$

where

$$\begin{aligned}
C_t^P &= e^{\phi p_t - t\nu_{p_c}} \\
C_t^\theta &= e^{\int_{t-2}^t g(t-s) dq_s - \nu_{q_c}}
\end{aligned}$$

and ν_{p_c} and ν_{q_c} are defined such that $\mathbb{E}(\frac{C_t^P}{C_s^P}) = \mathbb{E}(C_t^\theta) = 1$ for all $t \geq s$. Analogous to the case with log income growth over N years (equation 10) we get:

$$\Delta^N \log(\bar{c}_T) = \log\left(\frac{C_{T-1}^P}{C_{T-N}^P}\right) + \log\left(\int_{T-1}^T \frac{C_t^P}{C_{T-1}^P} C_t^\theta dt\right) - \log\left(\int_{T-N-1}^{T-N} \frac{C_t^P}{C_{T-N}^P} C_t^\theta dt\right) \quad (11)$$

Defining $\mathcal{C}_{T,N}^P$, $\mathcal{C}_{T,N}^1$ and $\mathcal{C}_{T,N}^2$ to be the three parts of the sum in equation 11 respectively, we have:

$$\mathcal{C}_{T,N}^P = \log\left(\frac{C_{T-1}^P}{C_{T-N}^P}\right)$$

$$\begin{aligned}
&= \phi \log\left(\frac{P_{T-1}}{P_{T-N}}\right) - (N-1)(\nu_{p_c} - \phi\nu_p) \\
\Rightarrow \text{Cov}(\mathcal{P}_{T,N}, \mathcal{C}_{T,N}^P) &= (N-1)\phi \text{Var}\left(\log\left(\frac{P_T}{P_{T-1}}\right)\right) \\
&= (N-1)\phi\sigma_P^2
\end{aligned}$$

and that:

$$\begin{aligned}
\mathcal{C}_{T,N}^1 &= \log\left(\int_{T-1}^T \frac{C_t^P}{C_{T-1}^P} C_t^\theta dt\right) \\
&= \log\left(\int_{T-1}^T \left(\frac{P_t}{P_{T-1}}\right)^\phi e^{-(t-(T-1))(\nu_{p_c}-\phi\nu_p)} C_t^\theta dt\right) \\
&\approx \log\left(\bar{C}_T^\theta\right) + \int_{T-1}^T \left(\left(\frac{P_t}{P_{T-1}}\right)^\phi e^{-(t-(T-1))(\nu_{p_c}-\phi\nu_p)} - 1\right) \frac{C_t^\theta}{\bar{C}_T^\theta} dt
\end{aligned}$$

where the steps taken in the approximation are the same as we did in the case of income.

$$\begin{aligned}
\text{Cov}\left(\mathcal{Q}_{T,N}^1, \mathcal{C}_{T,N}^1\right) &= \text{Cov}\left(\log\left(\bar{\theta}_T\right), \log\left(\bar{C}_T^\theta\right)\right) \\
&\quad + \mathbb{E}\left(\int_{T-1}^T \int_{T-1}^T \left(\frac{P_t}{P_{T-1}} - 1\right) \left(\left(\frac{P_s}{P_{T-1}}\right)^\phi e^{-(s-(T-1))(\nu_{p_c}-\phi\nu_p)} - 1\right) \frac{\theta_t}{\bar{\theta}_T} \frac{C_s^\theta}{\bar{C}_T^\theta} dt ds\right) \\
&= \text{Cov}\left(\log\left(\bar{\theta}_T\right), \log\left(\bar{C}_T^\theta\right)\right) \\
&\quad + \mathbb{E}\left(\int_{T-1}^T \int_{T-1}^T \left(\left(\frac{P_{\min(t,s)}}{P_{T-1}}\right)^{1+\phi} e^{-(\min(t,s)-(T-1))(\nu_{p_c}-\phi\nu_p)} - 1\right) \frac{\theta_t}{\bar{\theta}_T} \frac{C_s^\theta}{\bar{C}_T^\theta} dt ds\right) \\
&= \text{Cov}\left(\log\left(\bar{\theta}_T\right), \log\left(\bar{C}_T^\theta\right)\right) \\
&\quad + \int_{T-1}^T \int_{T-1}^T \mathbb{E}\left(\left(\frac{P_{\min(t,s)}}{P_{T-1}}\right)^{1+\phi} e^{-(\min(t,s)-(T-1))(\nu_{p_c}-\phi\nu_p)} - 1\right) dt ds \\
&\approx 0 \left\{ \begin{aligned} &+ \int_{T-1}^T \int_{T-1}^T \mathbb{E}\left(\left(\left(\frac{P_{\min(t,s)}}{P_{T-1}}\right)^{1+\phi} e^{-(\min(t,s)-(T-1))(\nu_{p_c}-\phi\nu_p)} - 1\right) \right. \\ &\quad \left. \times \left(\mathbb{E}(\hat{\theta}_t) + \mathbb{E}(\hat{C}_s^\theta) + \mathbb{E}(\hat{\theta}_t \hat{C}_s^\theta)\right)\right) dt ds \end{aligned} \right\} \\
&= \text{Cov}\left(\log\left(\bar{\theta}_T\right), \log\left(\bar{C}_T^\theta\right)\right) \\
&\quad + \int_0^1 \int_0^1 \mathbb{E}\left(P_{\min(t,s)}^{1+\phi} e^{-\min(t,s)(\nu_{p_c}-\phi\nu_p)} - 1\right) dt ds
\end{aligned}$$

where $\hat{C}_{t,T}^\theta = \frac{C_t^\theta - \bar{C}_T^\theta}{\bar{C}_T^\theta}$. We now assume that p_t has no jumps, and is therefore a Brownian motion. With this assumption, $\nu_p = \frac{1}{2}\sigma_P^2$ and $\nu_{p_c} = \frac{1}{2}\phi^2\sigma_P^2$ and $\mathbb{E}(P_t^{1+\phi}) =$

$e^{\frac{1}{2}t(1+\phi)^2\sigma_P^2 - \frac{1}{2}t(1+\phi)\sigma_P^2}$ so that:

$$\begin{aligned}
\text{Cov}\left(\mathcal{Q}_{T,N}^1, \mathcal{C}_{T,N}^1\right) &= \text{Cov}\left(\log\left(\bar{\theta}_T\right), \log\left(\bar{C}_T^\theta\right)\right) \\
&\quad + \int_0^1 \int_0^1 \left(e^{\frac{1}{2}\min(s,t)\sigma_P^2((1+\phi)^2 - (1+\phi) - \phi^2 + \phi)} - 1\right) dt ds \\
&= \text{Cov}\left(\log\left(\bar{\theta}_T\right), \log\left(\bar{C}_T^\theta\right)\right) \\
&\quad + \int_0^1 \int_0^1 \left(e^{\min(s,t)\phi\sigma_P^2} - 1\right) dt ds \\
&\approx \text{Cov}\left(\log\left(\bar{\theta}_T\right), \log\left(\bar{C}_T^\theta\right)\right) + \phi\sigma_P^2 \int_0^1 \int_0^1 \min(s,t) dt ds \\
&= \text{Cov}\left(\log\left(\bar{\theta}_T\right), \log\left(\bar{C}_T^\theta\right)\right) + \frac{1}{3}\phi\sigma_P^2
\end{aligned}$$

Similarly

$$\text{Cov}\left(\mathcal{Q}_{T,N}^2, \mathcal{C}_{T,N}^2\right) \approx \text{Cov}\left(\log\left(\bar{\theta}_T\right), \log\left(\bar{C}_T^\theta\right)\right) + \frac{1}{3}\phi\sigma_P^2$$

So that the covariance of income growth with consumption growth over N years is:

$$\text{Cov}\left(\Delta^N \log(\bar{y}_T), \Delta^N \log(\bar{c}_T)\right) = \left(N - \frac{1}{3}\right)\phi\sigma_P^2 + 2\text{Cov}(\tilde{y}, \tilde{c})$$

where $\tilde{y} = \log\left(\bar{\theta}_T\right)$ and $\tilde{c} = \log\left(\bar{C}_T^\theta\right)$

In this appendix we formally derive the key identification formulas. We start by defining permanent income. Let p_t for $t \in \mathbb{R}^+$ be a martingale process (possibly with jumps) with independent stationary increments and ν be such that $\mathbb{E}(e^{p_t - p_{t-1}}) = e^\nu$. Define permanent income as:

$$P_t = e^{p_t - t\nu}$$

Note that $\mathbb{E}\left(\frac{P_{t+s}}{P_t}\right) = 1$ for all $s \geq 0$. Define the variance of log permanent shocks to be:

$$\sigma_P^2 = \text{Var}\left(\log\left(\frac{P_{t+1}}{P_t}\right)\right) = \text{Var}(p_{t+1} - p_t)$$

We will assume changes in permanent income over a one year period are small enough such that:

$$\begin{aligned}
\text{Var}\left(\frac{P_{t+1}}{P_t}\right) &= \text{Var}\left(\frac{P_{t+1} - P_t}{P_t}\right) \\
&\approx \text{Var}\left(\log\left(1 + \frac{P_{t+1} - P_t}{P_t}\right)\right) \\
&= \text{Var}\left(\log\left(\frac{P_{t+1}}{P_t}\right)\right) = \sigma_P^2
\end{aligned} \tag{12}$$

For transitory shocks, we define an increasing stochastic process, Θ_t , which also has independent stationary increments. The increments in this process will define the transitory shocks. We set the expectation of increments, and the variance of the log of

an increment of length 1 as:

$$\begin{aligned}\mathbb{E}(\Theta_{t+s} - \Theta_t) &= s \\ \text{Var}\left(\log(\Theta_{t+1} - \Theta_t)\right) &= \sigma_{\Theta}^2\end{aligned}$$

Note that for this to be well defined, Θ_t must not only be increasing but also its increments are almost surely strictly positive (so that log of the increment is defined almost everywhere). Examples of such a stochastic process would be a gamma process, or a process that increases linearly with time (non-stochastically) but is also subject to positive shocks that arrive as a Poisson process. The stochastic part of this process has no Brownian motion component as this would necessarily lead to non-zero probability of a decreasing increment.

We will use these two processes to define an income process in discrete time with m intervals per period, and then look at the limit as $m \rightarrow \infty$. Define $\theta_{t,m}$ for $t \in \{\frac{1}{m}, \frac{2}{m}, \frac{3}{m}, \dots\}$ to be the increment of Θ_t from $t - \frac{1}{m}$ to t :

$$\theta_{t,m} = \Theta_t - \Theta_{t-\frac{1}{m}}$$

Income is defined for each period $t \in \{\frac{1}{m}, \frac{2}{m}, \frac{3}{m}, \dots\}$ as:

$$Y_{t,m} = P_t \theta_{t,m}$$

Therefore the underlying income process has a pure division into permanent and transitory shocks. Income is observed for $T \in \{1, 2, 3, \dots\}$ as the sum of income in each of the subperiods:

$$\bar{Y}_{T,m} = \sum_{i=0}^{m-1} P_{T-\frac{i}{m}} \theta_{T-\frac{i}{m},m}$$

Note that for $m = 1$ this is the same as the underlying income process, with permanent and transitory variance as defined above. We are interested in the log of observable income growth over N periods:

$$\begin{aligned}\Delta^N \bar{y}_{T,m} &= \log \bar{Y}_{T,m} - \log \bar{Y}_{T-N,m} \\ &= \log \left(\sum_{i=0}^{m-1} P_{T-\frac{i}{m}} \theta_{T-\frac{i}{m},m} \right) - \log \left(\sum_{i=0}^{m-1} P_{T-N-\frac{i}{m}} \theta_{T-N-\frac{i}{m},m} \right) \\ &= \log \left(P_{T-1} \sum_{i=0}^{m-1} \frac{P_{T-\frac{i}{m}}}{P_{T-1}} \theta_{T-\frac{i}{m},m} \right) - \log \left(P_{T-N} \sum_{i=0}^{m-1} \frac{P_{T-N-\frac{i}{m}}}{P_{T-N}} \theta_{T-N-\frac{i}{m},m} \right) \\ &= \log(P_{T-1}) - \log(P_{T-N}) \\ &\quad + \log \left(\sum_{i=0}^{m-1} \frac{P_{T-\frac{i}{m}}}{P_{T-1}} \theta_{T-\frac{i}{m},m} \right) - \log \left(\sum_{i=0}^{m-1} \frac{P_{T-N-\frac{i}{m}}}{P_{T-N}} \theta_{T-N-\frac{i}{m},m} \right)\end{aligned}$$

As P_t and Θ_t have independent increments, the covariance between each of the three parts of the sum above is 0. Therefore:

$$\text{Var}\left(\Delta^N \bar{y}_{T,m}\right) = \text{Var}\left(\log(P_{T-1}) - \log(P_{T-N})\right)$$

$$+ \text{Var} \left(\log \left(\sum_{i=0}^{m-1} \frac{P_{T-\frac{i}{m}}}{P_{T-1}} \theta_{T-\frac{i}{m},m} \right) \right) + \text{Var} \left(\log \left(\sum_{i=0}^{m-1} \frac{P_{T-N-\frac{i}{m}}}{P_{T-N}} \theta_{T-N-\frac{i}{m},m} \right) \right)$$

We will look at each of these three variances individually. Starting with the easiest:

$$\begin{aligned} \text{Var} \left(\log(P_{T-1}) - \log(P_{T-N}) \right) &= \text{Var} \left(\sum_{t=T-N+1}^{T-1} \log \left(\frac{P_t}{P_{t-1}} \right) \right) \\ &= (N-1) \sigma_P^2 \end{aligned}$$

Now consider the second variance. We begin by looking at the variable:

$$\begin{aligned} \log \left(\sum_{i=0}^{m-1} \frac{P_{T-\frac{i}{m}}}{P_{T-1}} \theta_{T-\frac{i}{m},m} \right) &= \log \left(\sum_{i=0}^{m-1} \theta_{T-\frac{i}{m},m} + \sum_{i=0}^{m-1} \left(\frac{P_{T-\frac{i}{m}}}{P_{T-1}} - 1 \right) \theta_{T-\frac{i}{m},m} \right) \\ &= \log \left(\Theta_T - \Theta_{T-1} \right) + \log \left(1 + \sum_{i=0}^{m-1} \left(\frac{P_{T-\frac{i}{m}}}{P_{T-1}} - 1 \right) \frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \right) \\ &\approx \log \left(\Theta_T - \Theta_{T-1} \right) + \sum_{i=0}^{m-1} \left(\frac{P_{T-\frac{i}{m}}}{P_{T-1}} - 1 \right) \frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \end{aligned}$$

Where the approximation comes from the fact that the shocks to permanent income in a one year period are small. Defining

$$\zeta_{t,m} = \frac{P_t}{P_{t-\frac{1}{m}}}$$

we have that

$$\begin{aligned} \text{Var} \left(\log \left(\sum_{i=0}^{m-1} \frac{P_{T-\frac{i}{m}}}{P_{T-1}} \theta_{T-\frac{i}{m},m} \right) \right) &\approx \sigma_\Theta^2 + \text{Var} \left(\sum_{i=0}^{m-1} \left(\prod_{j=i}^{m-1} \zeta_{T-\frac{j}{m}} - 1 \right) \frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \right) \\ &= \sigma_\Theta^2 + \mathbb{E} \left[\sum_{i=0}^{m-1} \left(\prod_{j=i}^{m-1} \zeta_{T-\frac{j}{m}} - 1 \right) \frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \right]^2 \\ &= \sigma_\Theta^2 + \mathbb{E} \left[\sum_{i=0}^{m-1} \left(\left(\prod_{j=i}^{m-1} \zeta_{T-\frac{j}{m}} - 1 \right)^2 \left(\frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \right)^2 \right. \right. \\ &\quad \left. \left. + 2 \sum_{k < i} \left(\prod_{j=k}^{m-1} \zeta_{T-\frac{j}{m}} - 1 \right) \left(\prod_{j=i}^{m-1} \zeta_{T-\frac{j}{m}} - 1 \right) \frac{\theta_{T-\frac{k}{m},m} \theta_{T-\frac{i}{m},m}}{\left(\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m} \right)^2} \right) \right] \\ &= \sigma_\Theta^2 + \frac{\sigma_P^2}{m} \sum_{i=0}^{m-1} \left(i \mathbb{E} \left(\frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \right)^2 + 2 \sum_{k < i} (m-1-i) \mathbb{E} \left(\frac{\theta_{T-\frac{k}{m},m} \theta_{T-\frac{i}{m},m}}{\left(\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m} \right)^2} \right) \right) \\ &= \sigma_\Theta^2 + \frac{\sigma_P^2}{m} \frac{m(m-1)}{2} \mathbb{E} \left(\frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \right)^2 \end{aligned}$$

$$\begin{aligned}
& + 2 \frac{\sigma_P^2}{m} \sum_{i=1}^{m-1} i(m-1-i) \mathbb{E} \left(\frac{\theta_{T-\frac{k}{m},m} \theta_{T-\frac{i}{m},m}}{\left(\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m} \right)^2} \right) \\
& = \sigma_\Theta^2 + \sigma_P^2 \frac{m-1}{2} \mathbb{E} \left(\frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \right)^2 \\
& \quad + \sigma_P^2 \left[(m-1)^2 - \frac{(m-1)(2m-1)}{3} \right] \mathbb{E} \left(\frac{\theta_{T-\frac{k}{m},m} \theta_{T-\frac{i}{m},m}}{\left(\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m} \right)^2} \right)
\end{aligned}$$

Note that:

$$\begin{aligned}
1 & = \mathbb{E} \left(\sum_{i=0}^{m-1} \frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \right)^2 \\
& = \sum_{i=0}^{m-1} \mathbb{E} \left(\frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \right)^2 + 2 \sum_{k < i} \mathbb{E} \left(\frac{\theta_{T-\frac{k}{m},m} \theta_{T-\frac{i}{m},m}}{\left(\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m} \right)^2} \right)
\end{aligned}$$

So that

$$\mathbb{E} \left(\frac{\theta_{T-\frac{k}{m},m} \theta_{T-\frac{i}{m},m}}{\left(\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m} \right)^2} \right) = \frac{1}{m(m-1)} - \frac{1}{m-1} \mathbb{E} \left(\frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \right)^2$$

This gives:

$$\begin{aligned}
\text{Var} \left(\log \left(\sum_{i=0}^{m-1} \frac{P_{T-\frac{i}{m}}}{P_{T-1}} \theta_{T-\frac{i}{m},m} \right) \right) & \approx \sigma_\Theta^2 + \text{Var} \left(\sum_{i=0}^{m-1} \left(\prod_{j=i}^{m-1} \zeta_{T-\frac{j}{m}} - 1 \right) \frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \right) \\
& \approx \sigma_\Theta^2 + \frac{m-2}{3m} \sigma_P^2 + \frac{m+1}{6} \mathbb{E} \left(\frac{\theta_{T-\frac{i}{m},m}}{\sum_{l=0}^{m-1} \theta_{T-\frac{l}{m},m}} \right)^2 \sigma_P^2 \\
& \rightarrow \sigma_\Theta^2 + \frac{1}{3} \sigma_P^2 \quad \text{as } m \rightarrow \infty
\end{aligned}$$

A very similar calculation shows that:

$$\text{Var} \left(\log \left(\sum_{i=0}^{m-1} \frac{P_{T-N-\frac{i}{m}}}{P_{T-N}} \theta_{T-N-\frac{i}{m},m} \right) \right) \rightarrow \sigma_\Theta^2 + \frac{1}{3} \sigma_P^2 \quad \text{as } m \rightarrow \infty$$

Putting this all together gives:

$$\text{Var} \left(\Delta^N \bar{y}_{T,m} \right) \rightarrow (N - \frac{1}{3}) \sigma_P^2 + 2\sigma_\Theta^2 \quad \text{as } m \rightarrow \infty$$

B Sample Selection

We choose to look at households whose head is between the ages of 30 and 55 in 2008. This is driven by the desire to remove households for who the assumption that most of

the income growth is unexpected. For the old and the young it is likely that individual households will have a lot of information about their income path that is not available to the econometrician (for example the year in which they plan to retire, or the fact that they are on a specific career track with set expectations of promotion and pay raises). We also want to remove households whose income volatility is increasing or decreasing sharply. Figures 10 and 11 show how our estimates of both income variance and MPX vary with age. The dots represent the point estimate for each age while the lines are the centered moving averages over the five nearest age groups. The solid black line shows the total variance of income growth over 1 year. It should not be surprising that income growth for households with heads in their 20's is highly volatile. This volatility plateaus around the age of 35 and stays at a constant level until the point of retirement at which point it temporarily grows before falling to an even lower level. We can see that while both transitory and permanent shocks to income are high early in life, permanent income shocks are particularly large while individuals find their place in the workforce. From the age of 30 to 55 both transitory and permanent shocks are approximately the same size and remarkably stable. At retirement shocks to permanent income rise, not surprisingly as retirement itself will be seen in the model as a shock, even as transitory income variance declines.

As the model assumes the variance to permanent and transitory shocks is constant in the observed period, interpretation of the numbers outside of the 30-55 age group needs to be treated with care. However, the figure clearly shows that within this age group the assumption of constant variance appears to be a reasonable one.

The dotted black line shows the variance of Δy assuming no persistence in the transitory component. The fact that this line is slightly above the empirical variance of Δy is consistent with some persistence in the transitory component of income, justifying our decision exclude growth over one and two years in our identification.

The level of both permanent and transitory shock variance for households aged 30 to 55 is approximately 0.0035, reflecting a standard deviation of 6%. Estimates using US data are significantly higher, especially for the transitory shock variance (for example Carroll and Samwick (1997) estimate 0.02 for permanent and 0.04 for transitory). This difference may be due to lower income inequality in Denmark, more progressive taxation and more generous unemployment insurance. The lower transitory variance will also be due to significantly reduced measurement error relative to the survey based US data.

C Danish Mortgage Market

D Distribution of Permanent and Transitory MPX by URE, NNP and Income




Figure 10 Permanent and Transitory Shock Variance by Age




Figure 11 MPX by Age




Figure 12 Mortgage Debt by Type




Figure 13 Mortgage Debt by Maturity




Figure 14 MPX Distribution by URE, NNP and Income